

 Navigation

Machine Learning Mastery

Making Developers Awesome at Machine Learning

[Click to Take the FREE GANs Crash-Course](#)

Search...



How to Develop a Pix2Pix GAN for Image-to-Image Translation

by Jason Brownlee on August 2, 2019 in [Generative Adversarial Networks](#)

[Tweet](#)

[Share](#)

[Share](#)

Last Updated on January 18, 2021

The Pix2Pix Generative Adversarial Network, or GAN, is an approach to training a deep convolutional neural network for image-to-image translation tasks.

The careful configuration of architecture as a type of image-conditional GAN allows for both the generation of large images compared to prior GAN models (e.g. such as 256×256 pixels) and the capability of performing well on a variety of different image-to-image translation tasks.

In this tutorial, you will discover how to develop a Pix2Pix generative adversarial network for image-to-image translation.

After completing this tutorial, you will know:

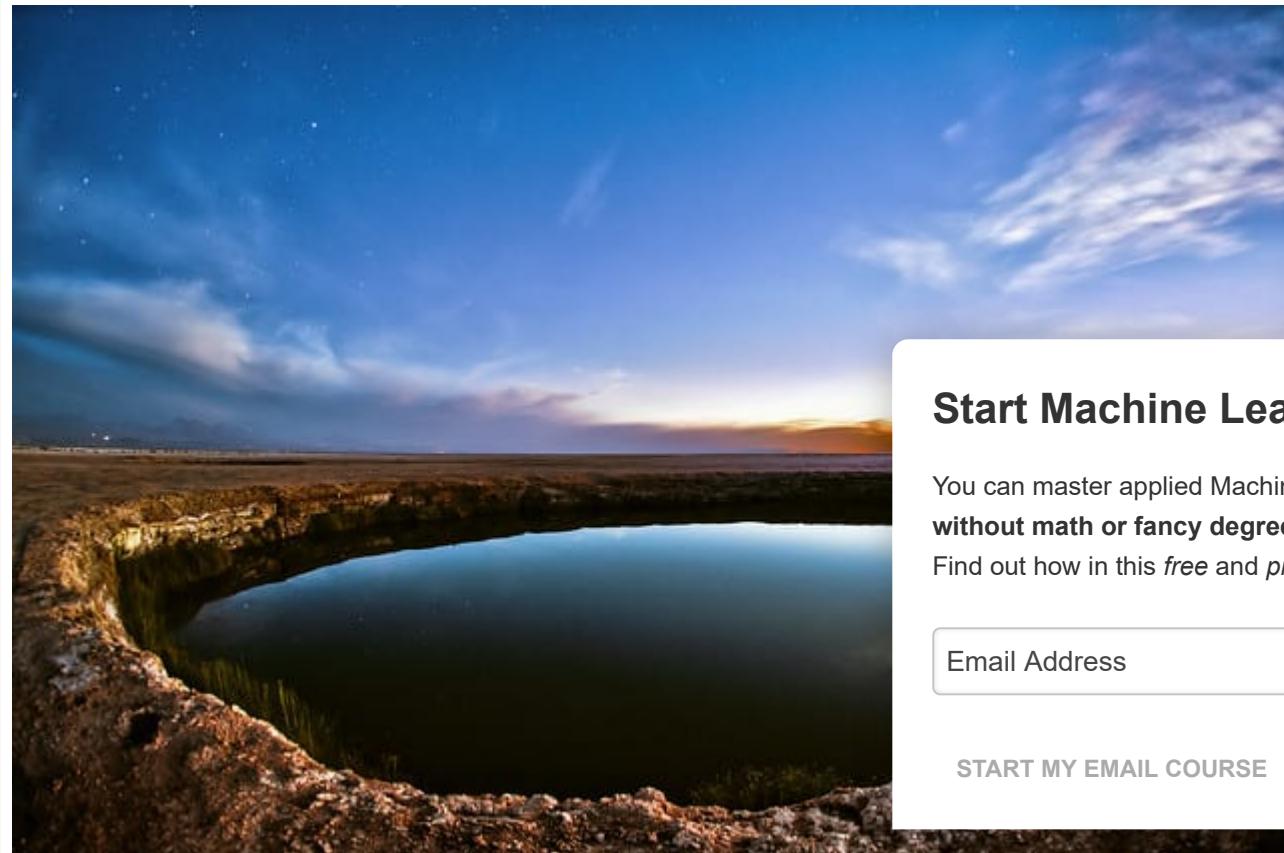
- How to load and prepare the satellite image to Google maps image-to-image translation dataset.
- How to develop a Pix2Pix model for translating satellite photographs to Google map images.
- How to use the final Pix2Pix generator model to translate ad hoc satellite images.

Kick-start your project with my new book [Generative Adversarial Networks with Python](#), including step-by-step tutorials and the [Python source code files](#) for all examples.

[Start Machine Learning](#)

Let's get started.

- **Updated Jan/2021:** Updated so layer freezing works with batch norm.



Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

How to Develop a Pix2Pix Generative Adversarial Network for Image-to-Image Translation

Photo by European Southern Observatory, some rights reserved.

Tutorial Overview

This tutorial is divided into five parts; they are:

1. What Is the Pix2Pix GAN?
2. Satellite to Map Image Translation Dataset
3. How to Develop and Train a Pix2Pix Model

[Start Machine Learning](#)

4. How to Translate Images With a Pix2Pix Model
5. How to Translate Google Maps to Satellite Images

What Is the Pix2Pix GAN?

Pix2Pix is a Generative Adversarial Network, or GAN, model designed for general purpose image-to-image translation.

The approach was presented by Phillip Isola, et al. in their 2016 paper titled “Image-to-Image Translation with Conditional Adversarial Networks” and presented at CVPR in 2017.

The GAN architecture is comprised of a generator model for outputting new plausible synthetic images, and a discriminator model that classifies images as real (from the dataset) or fake (generated). The discriminator model is updated directly, whereas the generator model is updated via the discriminator model. As such, the two models are trained simultaneously in a adversarial loop to better fool the discriminator and the discriminator seeks to better identify the counterfeit images.

The Pix2Pix model is a type of conditional GAN, or cGAN, where the generation of the output image is conditioned on a source image. The discriminator is provided both with a source image and the target image to distinguish between the plausible transformation of the source image.

The generator is trained via adversarial loss, which encourages the generator to generate images that look like the target image. The generator is also updated via L1 loss measured between the generated image and the expected output image. This allows the generator model to create plausible translations of the source image.

The Pix2Pix GAN has been demonstrated on a range of image-to-image translation tasks such as converting maps to satellite photographs, black and white photographs to color, and sketches of products to product photographs.

Now that we are familiar with the Pix2Pix GAN, let's prepare a dataset that we can use with image-to-image translation.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Want to Develop GANs from Scratch?

Take my free 7-day email crash course now (with sample code).

Click to sign-up and also get a free PDF Ebook version

Start Machine Learning

[Download Your FREE Mini-Course](#)

Satellite to Map Image Translation Dataset

In this tutorial, we will use the so-called “maps” dataset used in the Pix2Pix paper.

This is a dataset comprised of satellite images of New York and their corresponding Google maps pages. The image translation problem involves converting satellite photos to Google maps format, or the reverse, Google maps images to Satellite photos.

The dataset is provided on the pix2pix website and can be downloaded as a 255-megabyte file:

- [Download Maps Dataset \(maps.tar.gz\)](#)

Download the dataset and unzip it into your current working directory. This will create a directory structure like this:

```
1 maps
2   └── train
3     └── val
```

The train folder contains 1,097 images, whereas the validation dataset contains 1,099 images.

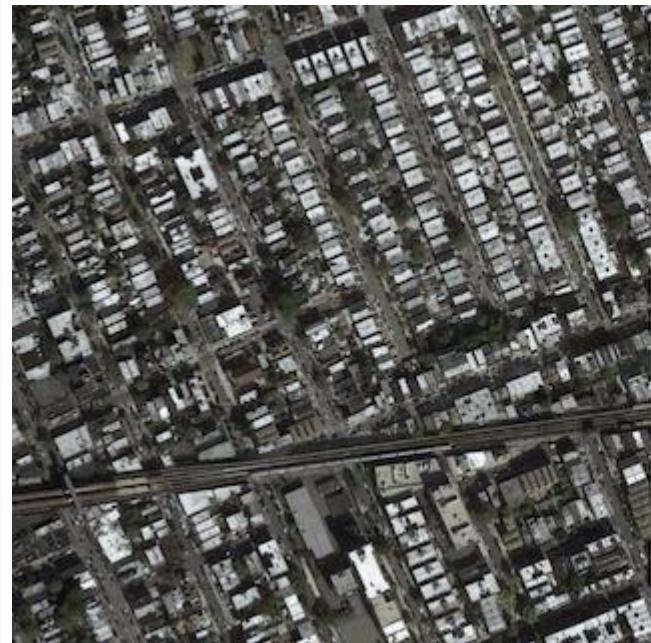
Images have a digit filename and are in JPEG format. Each image is 1,200 pixels wide and 675 pixels high, with the satellite image on the left and the Google maps image on the right.

Start Machine Learning ×

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

[Start Machine Learning](#)



Sample Image From the Maps Dataset Including Both Satellite and Map Elements

We can prepare this dataset for training a Pix2Pix GAN model in Keras. We will just work with the images as they are. The images will be loaded, rescaled, and split into the satellite and Google map elements. The result will be a list of arrays of images with a height of 256×256 pixels.

The `load_images()` function below implements this. It enumerates the list of images in a given directory, loads each image at 256×512 pixels, splits each image into satellite and map elements and returns an array of

```
1 # load all images in a directory into memory
2 def load_images(path, size=(256,512)):
3     src_list, tar_list = list(), list()
4     # enumerate filenames in directory, assume all are images
5     for filename in listdir(path):
6         # load and resize the image
7         pixels = load_img(path + filename, target_size=size)
8         # convert to numpy array
9         pixels = img_to_array(pixels)
10        # split into satellite and map
11        sat_img, map_img = pixels[:, :256], pixels[:, 256:]
12        src_list.append(sat_img)
13        tar_list.append(map_img)
14    return [asarray(src_list), asarray(tar_list)]
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

We can call this function with the path to the training dataset. Once loaded, we can save the prepared arrays to a new file in compressed format for later use.

The complete example is listed below.

```

1 # load, split and scale the maps dataset ready for training
2 from os import listdir
3 from numpy import asarray
4 from numpy import vstack
5 from keras.preprocessing.image import img_to_array
6 from keras.preprocessing.image import load_img
7 from numpy import savez_compressed
8
9 # load all images in a directory into memory
10 def load_images(path, size=(256,512)):
11     src_list, tar_list = list(), list()
12     # enumerate filenames in directory, assume all are images
13     for filename in listdir(path):
14         # load and resize the image
15         pixels = load_img(path + filename, target_size=size)
16         # convert to numpy array
17         pixels = img_to_array(pixels)
18         # split into satellite and map
19         sat_img, map_img = pixels[:, :256], pixels[:, 256:]
20         src_list.append(sat_img)
21         tar_list.append(map_img)
22     return [asarray(src_list), asarray(tar_list)]
23
24 # dataset path
25 path = 'maps/train/'
26 # load dataset
27 [src_images, tar_images] = load_images(path)
28 print('Loaded: ', src_images.shape, tar_images.shape)
29 # save as compressed numpy array
30 filename = 'maps_256.npz'
31 savez_compressed(filename, src_images, tar_images)
32 print('Saved dataset: ', filename)
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

Running the example loads all images in the training dataset, summarizes their shape to ensure the images were loaded correctly, then saves the arrays to a new file called *maps_256.npz* in compressed NumPy array format.

```

1 Loaded:  (1096, 256, 256, 3) (1096, 256, 256, 3)
2 Saved dataset:  maps_256.npz
```

This file can be loaded later via the `load()` NumPy function and retrieving each array in turn.

[Start Machine Learning](#)

We can then plot some images pairs to confirm the data has been handled correctly.

```
1 # load the prepared dataset
2 from numpy import load
3 from matplotlib import pyplot
4 # load the dataset
5 data = load('maps_256.npz')
6 src_images, tar_images = data['arr_0'], data['arr_1']
7 print('Loaded: ', src_images.shape, tar_images.shape)
8 # plot source images
9 n_samples = 3
10 for i in range(n_samples):
11     pyplot.subplot(2, n_samples, 1 + i)
12     pyplot.axis('off')
13     pyplot.imshow(src_images[i].astype('uint8'))
14 # plot target image
15 for i in range(n_samples):
16     pyplot.subplot(2, n_samples, 1 + n_samples + i)
17     pyplot.axis('off')
18     pyplot.imshow(tar_images[i].astype('uint8'))
19 pyplot.show()
```

Running this example loads the prepared dataset and summarizes the shape of each array. It creates a 2x3 grid of plots showing three pairs of satellite images and Google map images.

```
1 Loaded: (1096, 256, 256, 3) (1096, 256, 256, 3)
```

A plot of three image pairs is also created showing the satellite images on the top and Google map images on the bottom.

We can see that satellite images are quite complex and that although the Google map images are blurry, they still contain things like major roads, water, and parks.

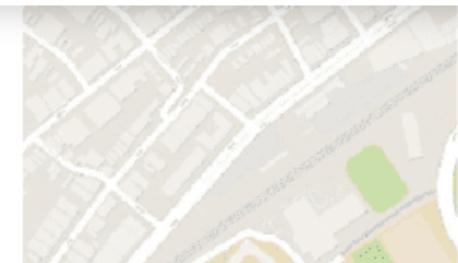
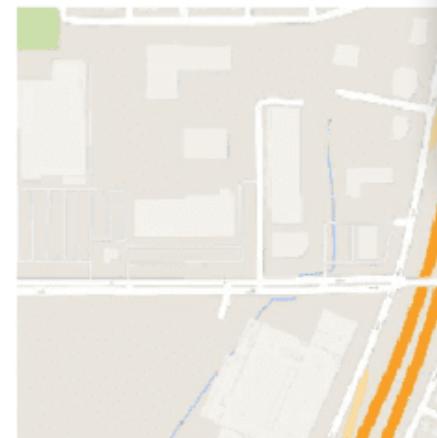
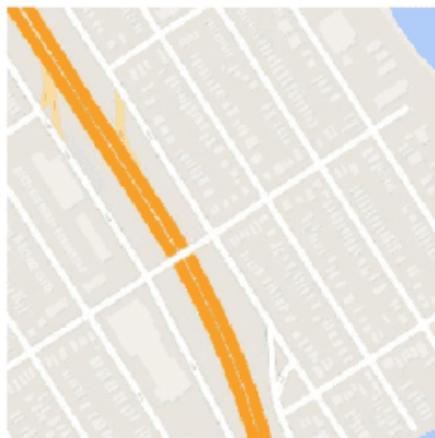
Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)



Start Machine Learning ×

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Start Machine Learning

Plot of Three Image Pairs Showing Satellite Images (top) and Google Map Images (bottom).

Now that we have prepared the dataset for image translation, we can develop our Pix2Pix GAN model.

How to Develop and Train a Pix2Pix Model

In this section, we will develop the Pix2Pix model for translating satellite photos to Google maps images.

The same model architecture and configuration described in the paper was used across a range of image translation tasks. This architecture is both described in the body of the paper, with additional detail in the appendix of the paper, and a [fully working implementation provided as open source with the Torch deep learning framework](#).

The implementation in this section will use the Keras deep learning framework based directly on the architecture implemented in the author's code base, designed to take and generate color images with the same dimensions.

The architecture is comprised of two models: the discriminator and the generator.

The discriminator is a deep convolutional neural network that performs image classification. It takes both the source image (e.g. satellite photo) and the target image (e.g. Google maps) and predicts whether target image is real or a fake translation of the source image.

The discriminator design is based on the effective receptive field of the model, which defines the number of pixels in the input image. This is called a PatchGAN model and is carefully designed so that the model maps to a 70×70 square or patch of the input image. The benefit of this approach is that the same model can be applied to input images of different sizes, e.g. larger or smaller than 256×256 pixels.

The output of the model depends on the size of the input image but may be one value or a square activation map of values. Each value is a probability for the likelihood that a patch in the input image is real. These values can be averaged to give an overall likelihood or classification score if needed.

The `define_discriminator()` function below implements the 70×70 PatchGAN discriminator model as per the design of the model in the paper. The model takes two input images that are concatenated together and predicts a patch output of predictions. The model is optimized using binary cross entropy, and a weighting is used so that updates to the model have half (0.5) the usual effect. The authors of Pix2Pix recommend this weighting of model updates to slow down changes to the discriminator, relative to the generator.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

```

1 # define the discriminator model
2 def define_discriminator(image_shape):
3     # weight initialization
4     init = RandomNormal(stddev=0.02)
5     # source image input
6     in_src_image = Input(shape=image_shape)
7     # target image input
8     in_target_image = Input(shape=image_shape)
9     # concatenate images channel-wise
10    merged = Concatenate()([in_src_image, in_target_image])
11    # C64
12    d = Conv2D(64, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(merged)
13    d = LeakyReLU(alpha=0.2)(d)
14    # C128
15    d = Conv2D(128, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d)
16    d = BatchNormalization()(d)
17    d = LeakyReLU(alpha=0.2)(d)
18    # C256
19    d = Conv2D(256, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d)
20    d = BatchNormalization()(d)
21    d = LeakyReLU(alpha=0.2)(d)
22    # C512
23    d = Conv2D(512, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d)
24    d = BatchNormalization()(d)
25    d = LeakyReLU(alpha=0.2)(d)
26    # second last output layer
27    d = Conv2D(512, (4,4), padding='same', kernel_initializer=init)(d)
28    d = BatchNormalization()(d)
29    d = LeakyReLU(alpha=0.2)(d)
30    # patch output
31    d = Conv2D(1, (4,4), padding='same', kernel_initializer=init)(d)
32    patch_out = Activation('sigmoid')(d)
33    # define model
34    model = Model([in_src_image, in_target_image], patch_out)
35    # compile model
36    opt = Adam(lr=0.0002, beta_1=0.5)
37    model.compile(loss='binary_crossentropy', optimizer=opt, loss_weights=[0.5])
38    return model

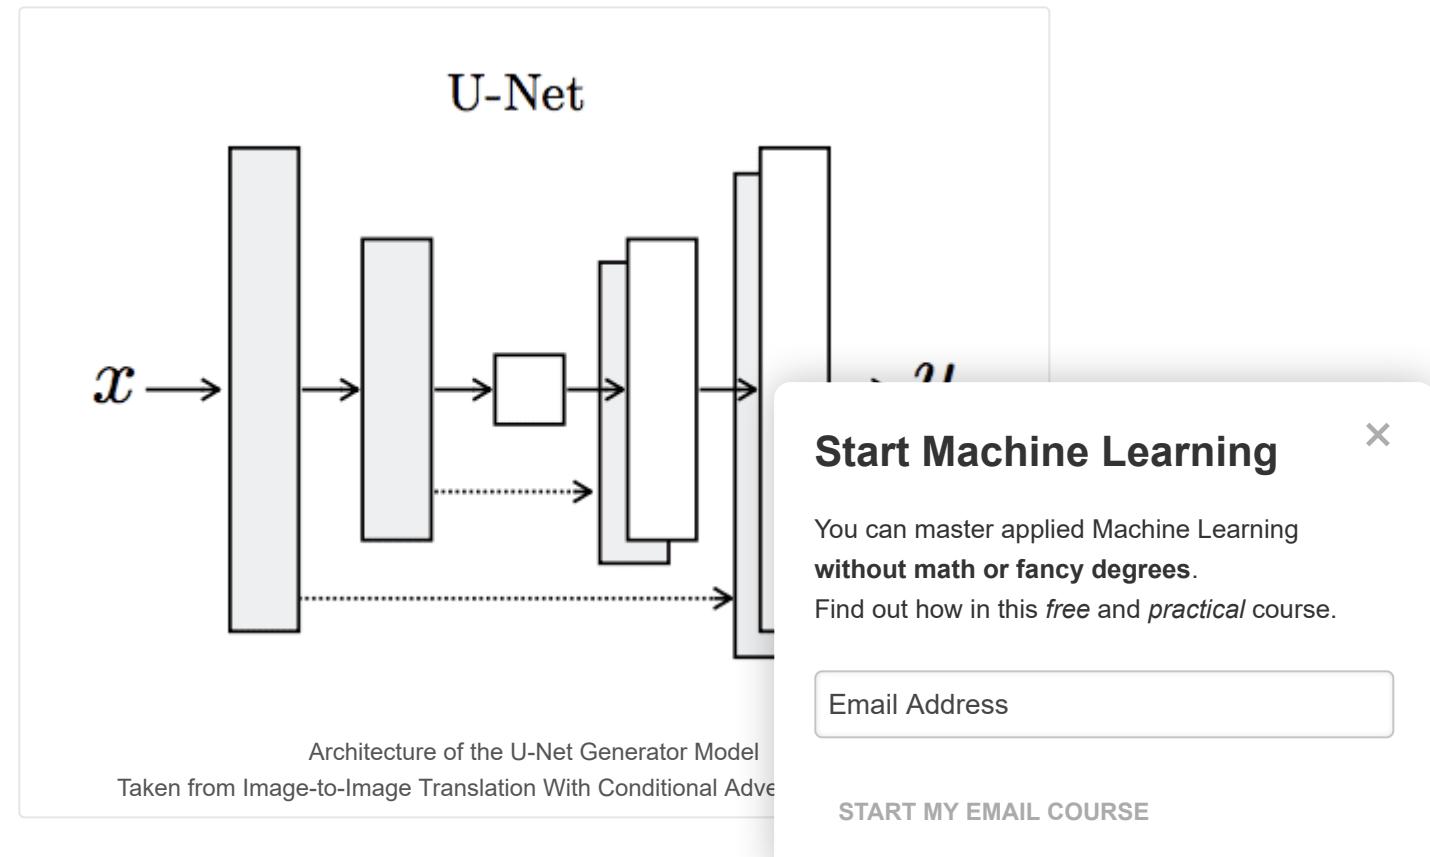
```

The generator model is more complex than the discriminator model.

The generator is an encoder-decoder model using a U-Net architecture. The model takes a source image (e.g. satellite photo) and generates a target image (e.g. Google maps image). It does this by first downsampling or encoding the input image down to a bottleneck layer, then upsampling or decoding the bottleneck representation to the size of the output image. The U-Net architecture means that skip-connections are added between the encoding layers and the corresponding decoding layers, forming a U-shape.

[Start Machine Learning](#)

The image below makes the skip-connections clear, showing how the first layer of the encoder is connected to the last layer of the decoder, and so on.



The encoder and decoder of the generator are comprised of standardized blocks of convolutional, batch normalization, dropout, and activation layers. This standardization means that we can develop helper functions to create each block of layers and call it repeatedly to build-up the encoder and decoder parts of the model.

The `define_generator()` function below implements the U-Net encoder-decoder generator model. It uses the `define_encoder_block()` helper function to create blocks of layers for the encoder and the `decoder_block()` function to create blocks of layers for the decoder. The tanh activation function is used in the output layer, meaning that pixel values in the generated image will be in the range [-1,1].

```
1 # define an encoder block
2 def define_encoder_block(layer_in, n_filters, batchnorm=True):
3     # weight initialization
4     init = RandomNormal(stddev=0.02)
5     # add downsampling layer
```

Start Machine Learning

```

6   g = Conv2D(n_filters, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(layer_in)
7   # conditionally add batch normalization
8   if batchnorm:
9     g = BatchNormalization()(g, training=True)
10  # leaky relu activation
11  g = LeakyReLU(alpha=0.2)(g)
12  return g
13
14 # define a decoder block
15 def decoder_block(layer_in, skip_in, n_filters, dropout=True):
16   # weight initialization
17   init = RandomNormal(stddev=0.02)
18   # add upsampling layer
19   g = Conv2DTranspose(n_filters, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(layer_in)
20   # add batch normalization
21   g = BatchNormalization()(g, training=True)
22   # conditionally add dropout
23   if dropout:
24     g = Dropout(0.5)(g, training=True)
25   # merge with skip connection
26   g = Concatenate()([g, skip_in])
27   # relu activation
28   g = Activation('relu')(g)
29   return g
30
31 # define the standalone generator model
32 def define_generator(image_shape=(256,256,3)):
33   # weight initialization
34   init = RandomNormal(stddev=0.02)
35   # image input
36   in_image = Input(shape=image_shape)
37   # encoder model
38   e1 = define_encoder_block(in_image, 64, batchnorm=False)
39   e2 = define_encoder_block(e1, 128)
40   e3 = define_encoder_block(e2, 256)
41   e4 = define_encoder_block(e3, 512)
42   e5 = define_encoder_block(e4, 512)
43   e6 = define_encoder_block(e5, 512)
44   e7 = define_encoder_block(e6, 512)
45   # bottleneck, no batch norm and relu
46   b = Conv2D(512, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(e7)
47   b = Activation('relu')(b)
48   # decoder model
49   d1 = decoder_block(b, e7, 512)
50   d2 = decoder_block(d1, e6, 512)
51   d3 = decoder_block(d2, e5, 512)
52   d4 = decoder_block(d3, e4, 512, dropout=False)
53   d5 = decoder_block(d4, e3, 256, dropout=False)
54   d6 = decoder_block(d5, e2, 128, dropout=False)

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)
[Start Machine Learning](#)

```

55     d7 = decoder_block(d6, e1, 64, dropout=False)
56     # output
57     g = Conv2DTranspose(3, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d7)
58     out_image = Activation('tanh')(g)
59     # define model
60     model = Model(in_image, out_image)
61     return model

```

The discriminator model is trained directly on real and generated images, whereas the generator model is not.

Instead, the generator model is trained via the discriminator model. It is updated to minimize the loss predicted by the discriminator for generated images marked as “real.” As such, it is encouraged to generate more real images. The generator is also updated to minimize the [L1 loss](#) or mean absolute error between the generated image and the target image.

The generator is updated via a weighted sum of both the adversarial loss and the L1 loss, with a weighting of 100 to 1 in favor of the L1 loss. This is to encourage the generator strongly to generate images that look like the target image, and not just plausible images in the target domain.

This can be achieved by defining a new logical model comprised of the weights in the existing generator and discriminator models. This logical or composite model involves stacking the generator on top of the discriminator model. The generator and discriminator are connected to each other, with the output of the generator connected to the input of the discriminator. The discriminator then predicts the likelihood that the generator was a real translation of the source image.

The discriminator is updated in a standalone manner, so the weights are reused in this context. The composite model is updated with two targets, one indicating that the generated images were real and the other indicating that they were not. These targets encourage updates in the generator toward generating more realistic images, and the executed real translation of the source image is used as the target for the output of the generator model (L1 loss).

The `define_gan()` function below implements this, taking the already-defined generator and discriminator models as arguments and using the [Keras functional API](#) to connect them together into a composite model. Both loss functions are specified for the two outputs of the model and the weights used for each are specified in the `loss_weights` argument to the `compile()` function.

```

1 # define the combined generator and discriminator model, for updating the generator
2 def define_gan(g_model, d_model, image_shape):
3     # make weights in the discriminator not trainable
4     for layer in d_model.layers:
5         if not isinstance(layer, BatchNormalization):
6             layer.trainable = False
7     # define the source image

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

```

8  in_src = Input(shape=image_shape)
9  # connect the source image to the generator input
10 gen_out = g_model(in_src)
11 # connect the source input and generator output to the discriminator input
12 dis_out = d_model([in_src, gen_out])
13 # src image as input, generated image and classification output
14 model = Model(in_src, [dis_out, gen_out])
15 # compile model
16 opt = Adam(lr=0.0002, beta_1=0.5)
17 model.compile(loss=['binary_crossentropy', 'mae'], optimizer=opt, loss_weights=[1,100])
18 return model

```

Next, we can load our paired images dataset in compressed NumPy array format.

This will return a list of two NumPy arrays: the first for source images and the second for corresponding target images.

```

1 # load and prepare training images
2 def load_real_samples(filename):
3     # load compressed arrays
4     data = load(filename)
5     # unpack arrays
6     X1, X2 = data['arr_0'], data['arr_1']
7     # scale from [0,255] to [-1,1]
8     X1 = (X1 - 127.5) / 127.5
9     X2 = (X2 - 127.5) / 127.5
10    return [X1, X2]

```

Training the discriminator will require batches of real and fake images.

The `generate_real_samples()` function below will prepare a batch of random pairs of images and a discriminator label of `class=1` to indicate they are real.

```

1 # select a batch of random samples, returns images and target
2 def generate_real_samples(dataset, n_samples, patch_shape):
3     # unpack dataset
4     trainA, trainB = dataset
5     # choose random instances
6     ix = randint(0, trainA.shape[0], n_samples)
7     # retrieve selected images
8     X1, X2 = trainA[ix], trainB[ix]
9     # generate 'real' class labels (1)
10    y = ones((n_samples, patch_shape, patch_shape, 1))
11    return [X1, X2], y

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

The `generate_fake_samples()` function below uses the generator model and a batch of real source images to generate an equivalent batch of target images for the discriminator.

These are returned with the label class-0 to indicate to the discriminator that they are fake.

```

1 # generate a batch of images, returns images and targets
2 def generate_fake_samples(g_model, samples, patch_shape):
3     # generate fake instance
4     X = g_model.predict(samples)
5     # create 'fake' class labels (0)
6     y = zeros((len(X), patch_shape, patch_shape, 1))
7     return X, y

```

Typically, GAN models do not converge; instead, an equilibrium is found between the generator and discriminator models. As such, we cannot easily judge when training should stop. Therefore, we can save the model and use it to generate images periodically during training, such as every 10 training epochs.

We can then review the generated images at the end of training and use the image quality to determine if training has converged.

The `summarize_performance()` function implements this, taking the generator model at a particular iteration number, in this case three, of translations of randomly selected images in the dataset. The generated images are then plotted as three rows of images and the plot saved to file. Additionally, the model is saved to disk to load later.

Both the image and model filenames include the training iteration number, allowing us to evaluate the model's performance over time.

```

1 # generate samples and save as a plot and save the model
2 def summarize_performance(step, g_model, dataset, n_samples=3):
3     # select a sample of input images
4     [X_realA, X_realB], _ = generate_real_samples(dataset, n_samples, 1)
5     # generate a batch of fake samples
6     X_fakeB, _ = generate_fake_samples(g_model, X_realA, 1)
7     # scale all pixels from [-1,1] to [0,1]
8     X_realA = (X_realA + 1) / 2.0
9     X_realB = (X_realB + 1) / 2.0
10    X_fakeB = (X_fakeB + 1) / 2.0
11    # plot real source images
12    for i in range(n_samples):
13        pyplot.subplot(3, n_samples, 1 + i)
14        pyplot.axis('off')
15        pyplot.imshow(X_realA[i])
16    # plot generated target image

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

```

17  for i in range(n_samples):
18      pyplot.subplot(3, n_samples, 1 + n_samples + i)
19      pyplot.axis('off')
20      pyplot.imshow(X_fakeB[i])
21  # plot real target image
22  for i in range(n_samples):
23      pyplot.subplot(3, n_samples, 1 + n_samples*2 + i)
24      pyplot.axis('off')
25      pyplot.imshow(X_realB[i])
26  # save plot to file
27  filename1 = 'plot_%06d.png' % (step+1)
28  pyplot.savefig(filename1)
29  pyplot.close()
30  # save the generator model
31  filename2 = 'model_%06d.h5' % (step+1)
32  g_model.save(filename2)
33  print('>Saved: %s and %s' % (filename1, filename2))

```

Finally, we can train the generator and discriminator models.

The `train()` function below implements this, taking the defined generator, discriminator, configuration, and dataset. The number of epochs is set at 100 to keep training times down, although 200 was used in the paper. The number of examples is set at 100 to keep training times down, although 200 was used in the paper.

Training involves a fixed number of training iterations. There are 1,097 images in the training dataset. A batch size of one means 1,097 training steps. The generator is trained for 10,970 training steps, and the model will run for 100 epochs, or a total of 109,700 training steps.

Each training step involves first selecting a batch of real examples, then using the generator to generate fake images, and finally updating the discriminator using the real source images. The discriminator is then updated with the batch of real images and then fake images.

Next, the generator model is updated providing the real source images as input and providing class labels of 1 (real) and the real target images as the expected outputs of the model required for calculating loss. The generator has two loss scores as well as the weighted sum score returned from the call to `train_on_batch()`. We are only interested in the weighted sum score (the first value returned) as it is used to update the model weights.

Finally, the loss for each update is reported to the console each training iteration and model performance is evaluated every 10 training epochs.

```

1 # train pix2pix model
2 def train(d_model, g_model, gan_model, dataset, n_epochs=100, n_batch=1):
3     # determine the output square shape of the discriminator

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

```

4 n_patch = d_model.output_shape[1]
5 # unpack dataset
6 trainA, trainB = dataset
7 # calculate the number of batches per training epoch
8 bat_per_epo = int(len(trainA) / n_batch)
9 # calculate the number of training iterations
10 n_steps = bat_per_epo * n_epochs
11 # manually enumerate epochs
12 for i in range(n_steps):
13     # select a batch of real samples
14     [X_realA, X_realB], y_real = generate_real_samples(dataset, n_batch, n_patch)
15     # generate a batch of fake samples
16     X_fakeB, y_fake = generate_fake_samples(g_model, X_realA, n_patch)
17     # update discriminator for real samples
18     d_loss1 = d_model.train_on_batch([X_realA, X_realB], y_real)
19     # update discriminator for generated samples
20     d_loss2 = d_model.train_on_batch([X_realA, X_fakeB], y_fake)
21     # update the generator
22     g_loss, _, _ = gan_model.train_on_batch(X_realA, [y_real, X_realB])
23     # summarize performance
24     print('>%d, d1[% .3f] d2[% .3f] g[% .3f]' % (i+1, d_loss1, d_loss2, g_loss))
25     # summarize model performance
26     if (i+1) % (bat_per_epo * 10) == 0:
27         summarize_performance(i, g_model, dataset)

```

Tying all of this together, the complete code example of training a Pix2Pix GAN to translate satellite images to maps is shown below.

```

1 # example of pix2pix gan for satellite to map image-to-image translation
2 from numpy import load
3 from numpy import zeros
4 from numpy import ones
5 from numpy.random import randint
6 from keras.optimizers import Adam
7 from keras.initializers import RandomNormal
8 from keras.models import Model
9 from keras.models import Input
10 from keras.layers import Conv2D
11 from keras.layers import Conv2DTranspose
12 from keras.layers import LeakyReLU
13 from keras.layers import Activation
14 from keras.layers import Concatenate
15 from keras.layers import Dropout
16 from keras.layers import BatchNormalization
17 from keras.layers import LeakyReLU
18 from matplotlib import pyplot
19
20 # define the discriminator model

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)
[Start Machine Learning](#)

```

21 def define_discriminator(image_shape):
22     # weight initialization
23     init = RandomNormal(stddev=0.02)
24     # source image input
25     in_src_image = Input(shape=image_shape)
26     # target image input
27     in_target_image = Input(shape=image_shape)
28     # concatenate images channel-wise
29     merged = Concatenate()([in_src_image, in_target_image])
30     # C64
31     d = Conv2D(64, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(merged)
32     d = LeakyReLU(alpha=0.2)(d)
33     # C128
34     d = Conv2D(128, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d)
35     d = BatchNormalization()(d)
36     d = LeakyReLU(alpha=0.2)(d)
37     # C256
38     d = Conv2D(256, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d)
39     d = BatchNormalization()(d)
40     d = LeakyReLU(alpha=0.2)(d)
41     # C512
42     d = Conv2D(512, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d)
43     d = BatchNormalization()(d)
44     d = LeakyReLU(alpha=0.2)(d)
45     # second last output layer
46     d = Conv2D(512, (4,4), padding='same', kernel_initializer=init)(d)
47     d = BatchNormalization()(d)
48     d = LeakyReLU(alpha=0.2)(d)
49     # patch output
50     d = Conv2D(1, (4,4), padding='same', kernel_initializer=init)(d)
51     patch_out = Activation('sigmoid')(d)
52     # define model
53     model = Model([in_src_image, in_target_image], patch_out)
54     # compile model
55     opt = Adam(lr=0.0002, beta_1=0.5)
56     model.compile(loss='binary_crossentropy', optimizer=opt, loss_weights=[0.5])
57     return model
58
59 # define an encoder block
60 def define_encoder_block(layer_in, n_filters, batchnorm=True):
61     # weight initialization
62     init = RandomNormal(stddev=0.02)
63     # add downsampling layer
64     g = Conv2D(n_filters, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(layer_in)
65     # conditionally add batch normalization
66     if batchnorm:
67         g = BatchNormalization()(g, training=True)
68     # leaky relu activation
69     g = LeakyReLU(alpha=0.2)(g)

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

```

70     return g
71
72 # define a decoder block
73 def decoder_block(layer_in, skip_in, n_filters, dropout=True):
74     # weight initialization
75     init = RandomNormal(stddev=0.02)
76     # add upsampling layer
77     g = Conv2DTranspose(n_filters, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(layer_in)
78     # add batch normalization
79     g = BatchNormalization()(g, training=True)
80     # conditionally add dropout
81     if dropout:
82         g = Dropout(0.5)(g, training=True)
83     # merge with skip connection
84     g = Concatenate()([g, skip_in])
85     # relu activation
86     g = Activation('relu')(g)
87     return g
88
89 # define the standalone generator model
90 def define_generator(image_shape=(256,256,3)):
91     # weight initialization
92     init = RandomNormal(stddev=0.02)
93     # image input
94     in_image = Input(shape=image_shape)
95     # encoder model
96     e1 = define_encoder_block(in_image, 64, batchnorm=False)
97     e2 = define_encoder_block(e1, 128)
98     e3 = define_encoder_block(e2, 256)
99     e4 = define_encoder_block(e3, 512)
100    e5 = define_encoder_block(e4, 512)
101    e6 = define_encoder_block(e5, 512)
102    e7 = define_encoder_block(e6, 512)
103    # bottleneck, no batch norm and relu
104    b = Conv2D(512, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(e7)
105    b = Activation('relu')(b)
106    # decoder model
107    d1 = decoder_block(b, e7, 512)
108    d2 = decoder_block(d1, e6, 512)
109    d3 = decoder_block(d2, e5, 512)
110    d4 = decoder_block(d3, e4, 512, dropout=False)
111    d5 = decoder_block(d4, e3, 256, dropout=False)
112    d6 = decoder_block(d5, e2, 128, dropout=False)
113    d7 = decoder_block(d6, e1, 64, dropout=False)
114    # output
115    g = Conv2DTranspose(3, (4,4), strides=(2,2), padding='same', kernel_initializer=init)(d7)
116    out_image = Activation('tanh')(g)
117    # define model
118    model = Model(in_image, out_image)

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)
[Start Machine Learning](#)

```

119     return model
120
121 # define the combined generator and discriminator model, for updating the generator
122 def define_gan(g_model, d_model, image_shape):
123     # make weights in the discriminator not trainable
124     for layer in d_model.layers:
125         if not isinstance(layer, BatchNormalization):
126             layer.trainable = False
127
128     # define the source image
129     in_src = Input(shape=image_shape)
130     # connect the source image to the generator input
131     gen_out = g_model(in_src)
132     # connect the source input and generator output to the discriminator input
133     dis_out = d_model([in_src, gen_out])
134     # src image as input, generated image and classification output
135     model = Model(in_src, [dis_out, gen_out])
136     # compile model
137     opt = Adam(lr=0.0002, beta_1=0.5)
138     model.compile(loss=['binary_crossentropy', 'mae'], optimizer=opt, loss_weights=[1, 100])
139     return model
140
141 # load and prepare training images
142 def load_real_samples(filename):
143     # load compressed arrays
144     data = load(filename)
145     # unpack arrays
146     X1, X2 = data['arr_0'], data['arr_1']
147     # scale from [0,255] to [-1,1]
148     X1 = (X1 - 127.5) / 127.5
149     X2 = (X2 - 127.5) / 127.5
150     return [X1, X2]
151
152 # select a batch of random samples, returns images and target
153 def generate_real_samples(dataset, n_samples, patch_shape):
154     # unpack dataset
155     trainA, trainB = dataset
156     # choose random instances
157     ix = randint(0, trainA.shape[0], n_samples)
158     # retrieve selected images
159     X1, X2 = trainA[ix], trainB[ix]
160     # generate 'real' class labels (1)
161     y = ones((n_samples, patch_shape, patch_shape, 1))
162     return [X1, X2], y
163
164 # generate a batch of images, returns images and targets
165 def generate_fake_samples(g_model, samples, patch_shape):
166     # generate fake instance
167     X = g_model.predict(samples)
168     # create 'fake' class labels (0)

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)
[Start Machine Learning](#)

```

168     y = zeros((len(X), patch_shape, patch_shape, 1))
169     return X, y
170
171 # generate samples and save as a plot and save the model
172 def summarize_performance(step, g_model, dataset, n_samples=3):
173     # select a sample of input images
174     [X_realA, X_realB], _ = generate_real_samples(dataset, n_samples, 1)
175     # generate a batch of fake samples
176     X_fakeB, _ = generate_fake_samples(g_model, X_realA, 1)
177     # scale all pixels from [-1,1] to [0,1]
178     X_realA = (X_realA + 1) / 2.0
179     X_realB = (X_realB + 1) / 2.0
180     X_fakeB = (X_fakeB + 1) / 2.0
181     # plot real source images
182     for i in range(n_samples):
183         pyplot.subplot(3, n_samples, 1 + i)
184         pyplot.axis('off')
185         pyplot.imshow(X_realA[i])
186     # plot generated target image
187     for i in range(n_samples):
188         pyplot.subplot(3, n_samples, 1 + n_samples + i)
189         pyplot.axis('off')
190         pyplot.imshow(X_fakeB[i])
191     # plot real target image
192     for i in range(n_samples):
193         pyplot.subplot(3, n_samples, 1 + n_samples*2 + i)
194         pyplot.axis('off')
195         pyplot.imshow(X_realB[i])
196     # save plot to file
197     filename1 = 'plot_%06d.png' % (step+1)
198     pyplot.savefig(filename1)
199     pyplot.close()
200     # save the generator model
201     filename2 = 'model_%06d.h5' % (step+1)
202     g_model.save(filename2)
203     print('>Saved: %s and %s' % (filename1, filename2))
204
205 # train pix2pix models
206 def train(d_model, g_model, gan_model, dataset, n_epochs=100, n_batch=1):
207     # determine the output square shape of the discriminator
208     n_patch = d_model.output_shape[1]
209     # unpack dataset
210     trainA, trainB = dataset
211     # calculate the number of batches per training epoch
212     bat_per_epo = int(len(trainA) / n_batch)
213     # calculate the number of training iterations
214     n_steps = bat_per_epo * n_epochs
215     # manually enumerate epochs
216     for i in range(n_steps):

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)
[Start Machine Learning](#)

```

217 # select a batch of real samples
218 [X_realA, X_realB], y_real = generate_real_samples(dataset, n_batch, n_patch)
219 # generate a batch of fake samples
220 X_fakeB, y_fake = generate_fake_samples(g_model, X_realA, n_patch)
221 # update discriminator for real samples
222 d_loss1 = d_model.train_on_batch([X_realA, X_realB], y_real)
223 # update discriminator for generated samples
224 d_loss2 = d_model.train_on_batch([X_realA, X_fakeB], y_fake)
225 # update the generator
226 g_loss, _, _ = gan_model.train_on_batch(X_realA, [y_real, X_realB])
227 # summarize performance
228 print('>%d, d1[%3f] d2[%3f] g[%3f]' % (i+1, d_loss1, d_loss2, g_loss))
229 # summarize model performance
230 if (i+1) % (bat_per_epo * 10) == 0:
231     summarize_performance(i, g_model, dataset)
232
233 # load image data
234 dataset = load_real_samples('maps_256.npz')
235 print('Loaded', dataset[0].shape, dataset[1].shape)
236 # define input shape based on the loaded dataset
237 image_shape = dataset[0].shape[1:]
238 # define the models
239 d_model = define_discriminator(image_shape)
240 g_model = define_generator(image_shape)
241 # define the composite model
242 gan_model = define_gan(g_model, d_model, image_shape)
243 # train model
244 train(d_model, g_model, gan_model, dataset)

```

The example can be run on CPU hardware, although GPU hardware is recommended.

The example might take about two hours to run on modern GPU hardware.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

The loss is reported each training iteration, including the discriminator loss on real examples (d1), discriminator loss on generated or fake examples (d2), and generator loss, which is a weighted average of adversarial and L1 loss (g).

If loss for the discriminator goes to zero and stays there for a long time, consider re-starting the training run as it is an example of a training failure.

1 >1, d1[0.566] d2[0.520] g[82.266]
 2 >2, d1[0.469] d2[0.484] g[66.813]

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

```
3 >3, d1[0.428] d2[0.477] g[79.520]
4 >4, d1[0.362] d2[0.405] g[78.143]
5 >5, d1[0.416] d2[0.406] g[72.452]
6 ...
7 >109596, d1[0.303] d2[0.006] g[5.792]
8 >109597, d1[0.001] d2[1.127] g[14.343]
9 >109598, d1[0.000] d2[0.381] g[11.851]
10 >109599, d1[1.289] d2[0.547] g[6.901]
11 >109600, d1[0.437] d2[0.005] g[10.460]
12 >Saved: plot_109600.png and model_109600.h5
```

Models are saved every 10 epochs and saved to a file with the training iteration number. Additionally, images are generated every 10 epochs and compared to the expected target images. These plots can be assessed at the end of the run and used to select a final generator model based on generated image quality.

At the end of the run, will you will have 10 saved model files and 10 plots of generated images.

After the first 10 epochs, map images are generated that look plausible, although the lines contain some blurring. Nevertheless, large structures are in the right places with mostly the

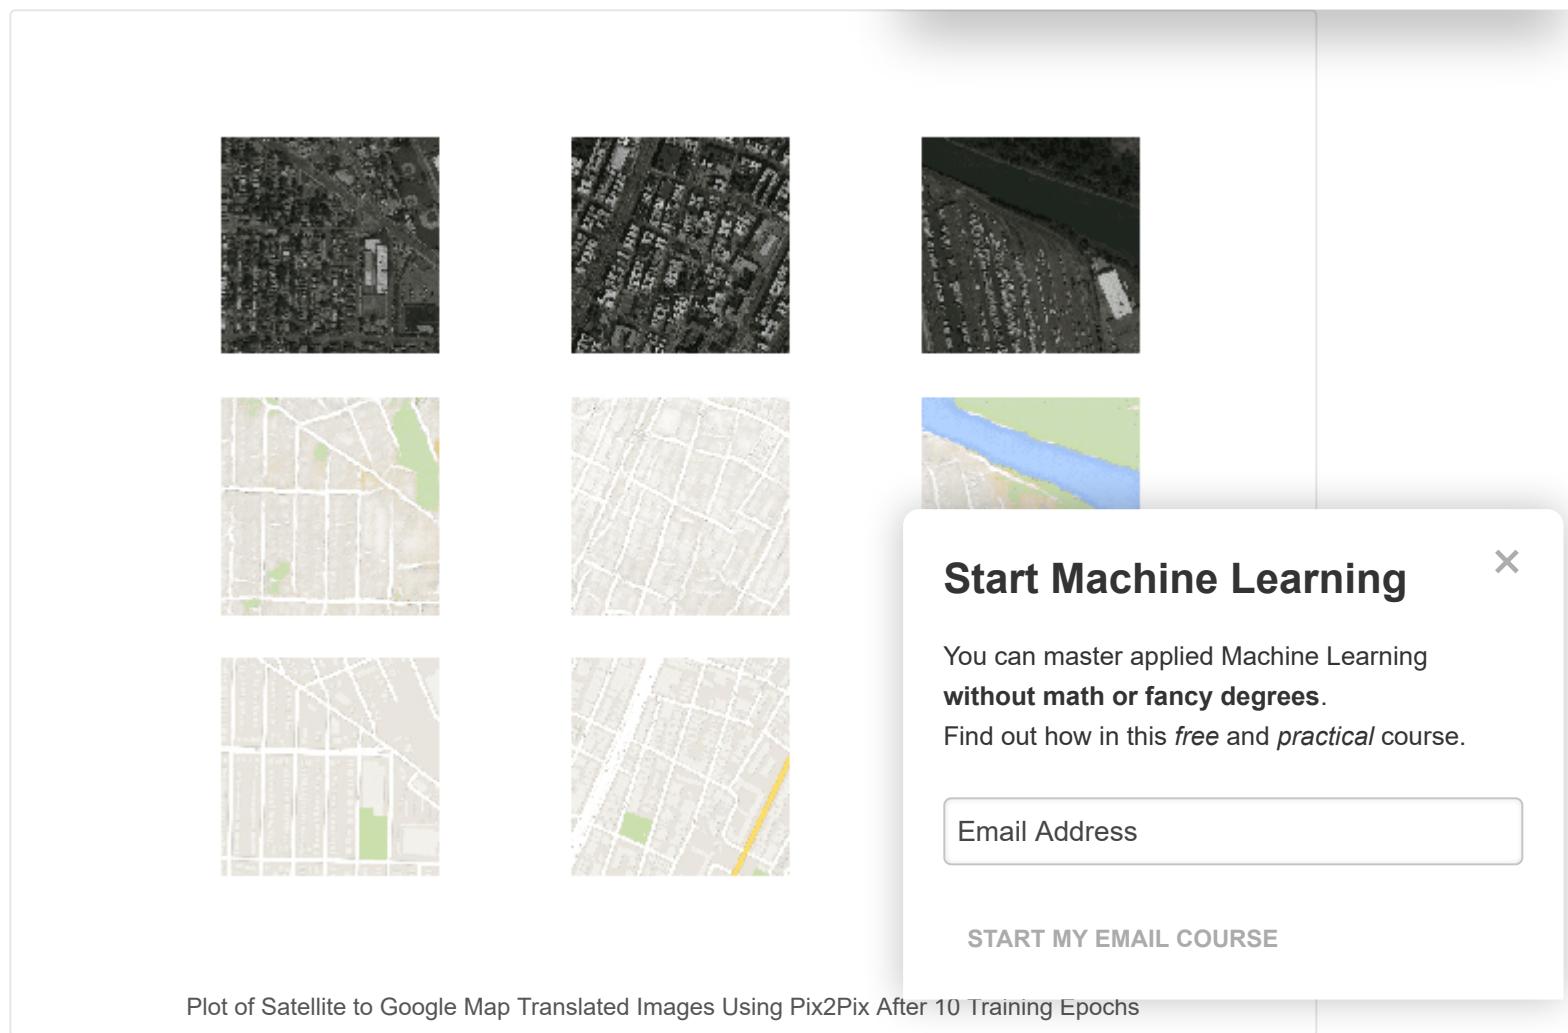
Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

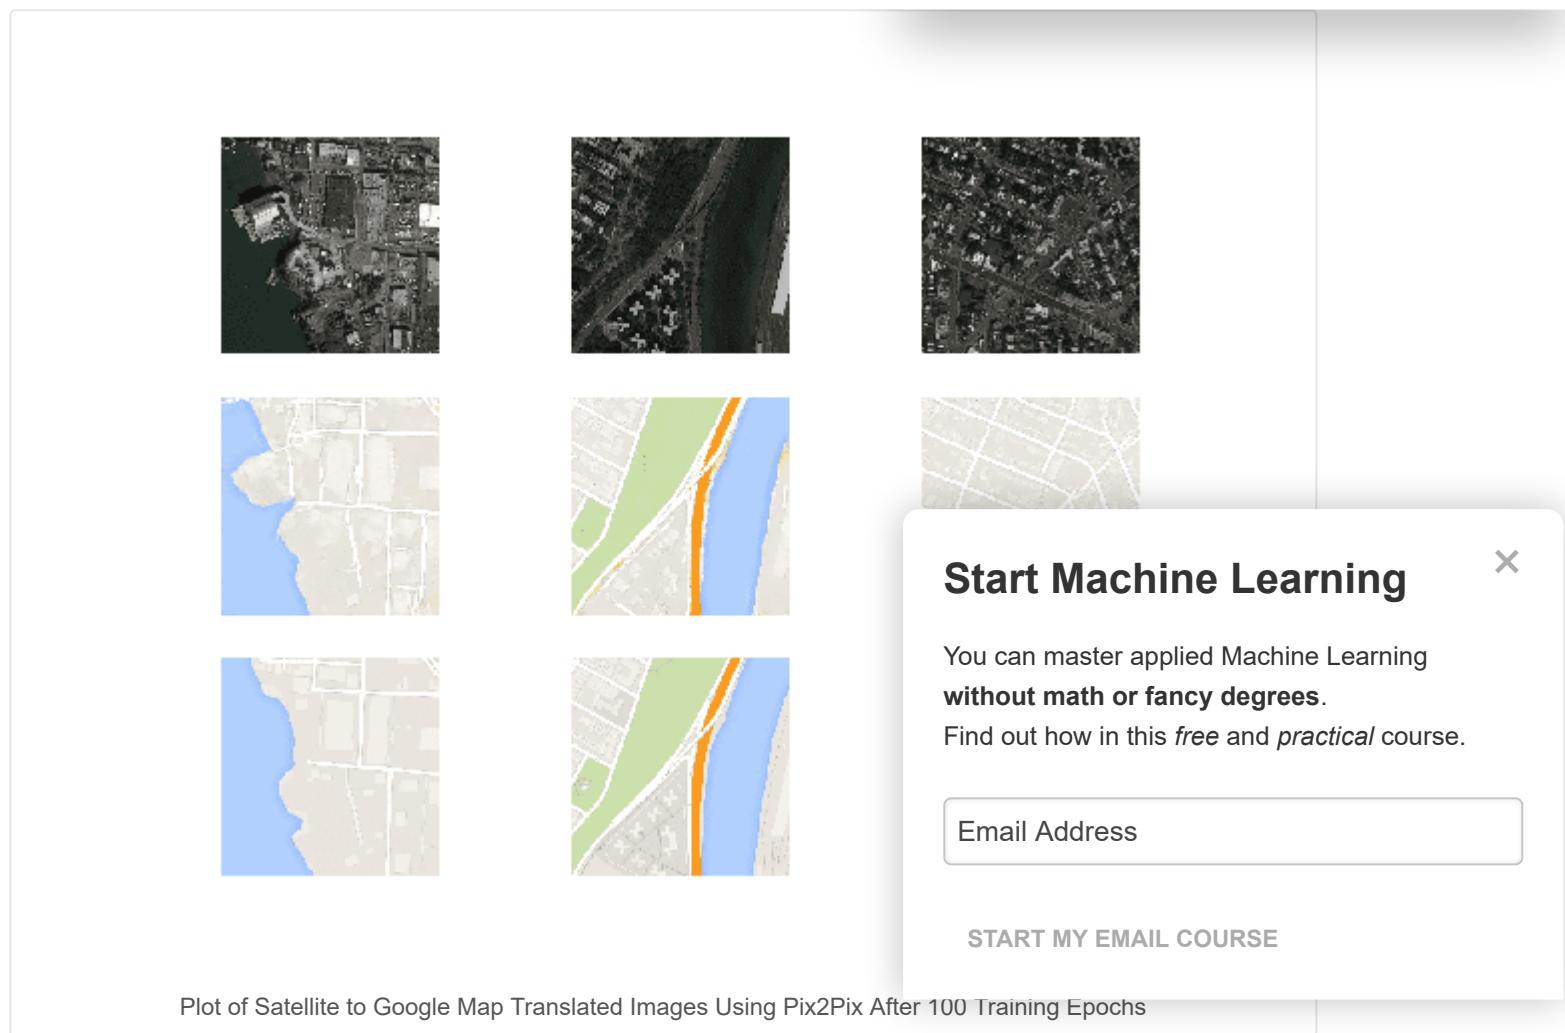
[Start Machine Learning](#)



Generated images after about 50 training epochs begin to look very realistic, at least to me, and quality appears to remain good for the remainder of the training process.

Note the first generated image example below (right column, middle row) that includes more useful detail than the real Google map image.

[Start Machine Learning](#)



Now that we have developed and trained the Pix2Pix model, we can explore how they can be used in a standalone manner.

How to Translate Images With a Pix2Pix Model

Training the Pix2Pix model results in many saved models and samples of generated images for each.

More training epochs does not necessarily mean a better quality model. Therefore, we can choose a model based on the quality of the generated images and use it to perform ad hoc image-to-image translation.

[Start Machine Learning](#)

In this case, we will use the model saved at the end of the run, e.g. after 100 epochs or 109,600 training iterations.

A good starting point is to load the model and use it to make ad hoc translations of source images in the training dataset.

First, we can load the training dataset. We can use the same function named *load_real_samples()* for loading the dataset as was used when training the model.

```

1 # load and prepare training images
2 def load_real_samples(filename):
3     # load compressed ararys
4     data = load(filename)
5     # unpack arrays
6     X1, X2 = data['arr_0'], data['arr_1']
7     # scale from [0,255] to [-1,1]
8     X1 = (X1 - 127.5) / 127.5
9     X2 = (X2 - 127.5) / 127.5
10    return [X1, X2]
```

This function can be called as follows:

```

1 ...
2 # load dataset
3 [X1, X2] = load_real_samples('maps_256.npz')
4 print('Loaded', X1.shape, X2.shape)
```

Next, we can load the saved Keras model.

```

1 ...
2 # load model
3 model = load_model('model_109600.h5')
```

Next, we can choose a random image pair from the training dataset to use as an example.

```

1 ...
2 # select random example
3 ix = randint(0, len(X1), 1)
4 src_image, tar_image = X1[ix], X2[ix]
```

We can provide the source satellite image as input to the model and use it to predict a Google map image.

```

1 ...
2 # generate image from source
3 gen_image = model.predict(src_image)
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

Finally, we can plot the source, generated image, and the expected target image.

The *plot_images()* function below implements this, providing a nice title above each image.

```

1 # plot source, generated and target images
2 def plot_images(src_img, gen_img, tar_img):
3     images = vstack((src_img, gen_img, tar_img))
4     # scale from [-1,1] to [0,1]
5     images = (images + 1) / 2.0
6     titles = ['Source', 'Generated', 'Expected']
7     # plot images row by row
8     for i in range(len(images)):
9         # define subplot
10        pyplot.subplot(1, 3, 1 + i)
11        # turn off axis
12        pyplot.axis('off')
13        # plot raw pixel data
14        pyplot.imshow(images[i])
15        # show title
16        pyplot.title(titles[i])
17    pyplot.show()

```

This function can be called with each of our source, generated, and target images.

```

1 ...
2 # plot all three images
3 plot_images(src_image, gen_image, tar_image)

```

Tying all of this together, the complete example of performing an ad hoc image-to-image translation is listed below.

```

1 # example of loading a pix2pix model and using it for image to image translation
2 from keras.models import load_model
3 from numpy import load
4 from numpy import vstack
5 from matplotlib import pyplot
6 from numpy.random import randint
7
8 # load and prepare training images
9 def load_real_samples(filename):
10    # load compressed arrays
11    data = load(filename)
12    # unpack arrays
13    X1, X2 = data['arr_0'], data['arr_1']
14    # scale from [0,255] to [-1,1]
15    X1 = (X1 - 127.5) / 127.5

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)
[Start Machine Learning](#)

```

16     X2 = (X2 - 127.5) / 127.5
17     return [X1, X2]
18
19 # plot source, generated and target images
20 def plot_images(src_img, gen_img, tar_img):
21     images = vstack((src_img, gen_img, tar_img))
22     # scale from [-1,1] to [0,1]
23     images = (images + 1) / 2.0
24     titles = ['Source', 'Generated', 'Expected']
25     # plot images row by row
26     for i in range(len(images)):
27         # define subplot
28         pyplot.subplot(1, 3, 1 + i)
29         # turn off axis
30         pyplot.axis('off')
31         # plot raw pixel data
32         pyplot.imshow(images[i])
33         # show title
34         pyplot.title(titles[i])
35     pyplot.show()
36
37 # load dataset
38 [X1, X2] = load_real_samples('maps_256.npz')
39 print('Loaded', X1.shape, X2.shape)
40 # load model
41 model = load_model('model_109600.h5')
42 # select random example
43 ix = randint(0, len(X1), 1)
44 src_image, tar_image = X1[ix], X2[ix]
45 # generate image from source
46 gen_image = model.predict(src_image)
47 # plot all three images
48 plot_images(src_image, gen_image, tar_image)

```

Running the example will select a random image from the training dataset, translate it to a Google map, and plot the result compared to the expected image.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

In this case, we can see that the generated image captures large roads with orange and yellow as well as green park areas. The generated image is not perfect but is very close to the expected image.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

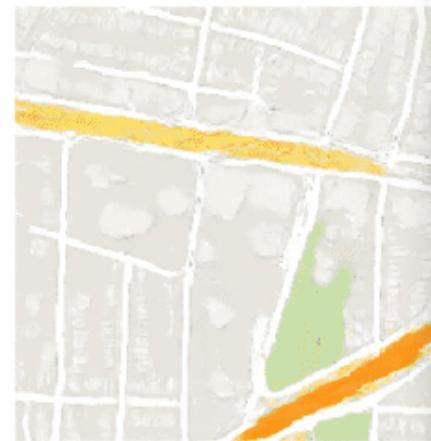
[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

Source



Generated



Expected

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

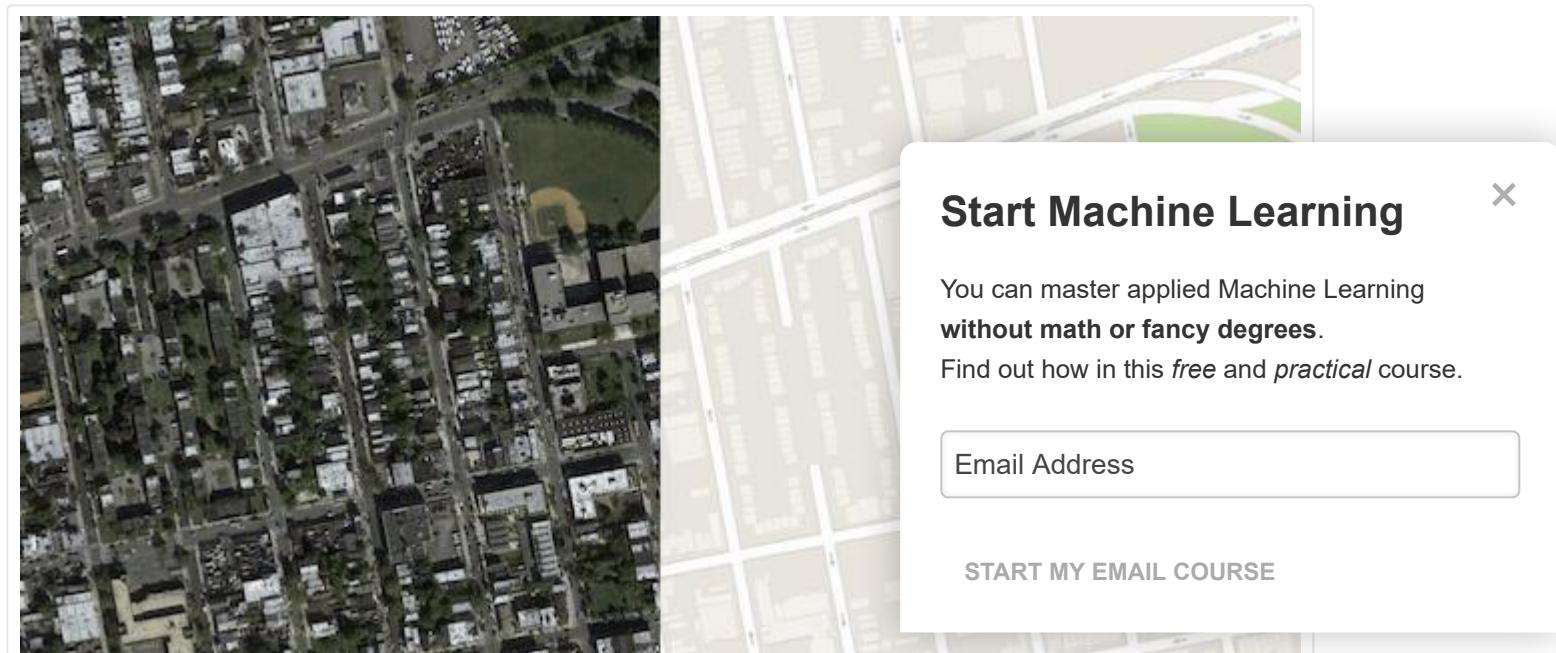
START MY EMAIL COURSE

Start Machine Learning

We may also want to use the model to translate a given standalone image.

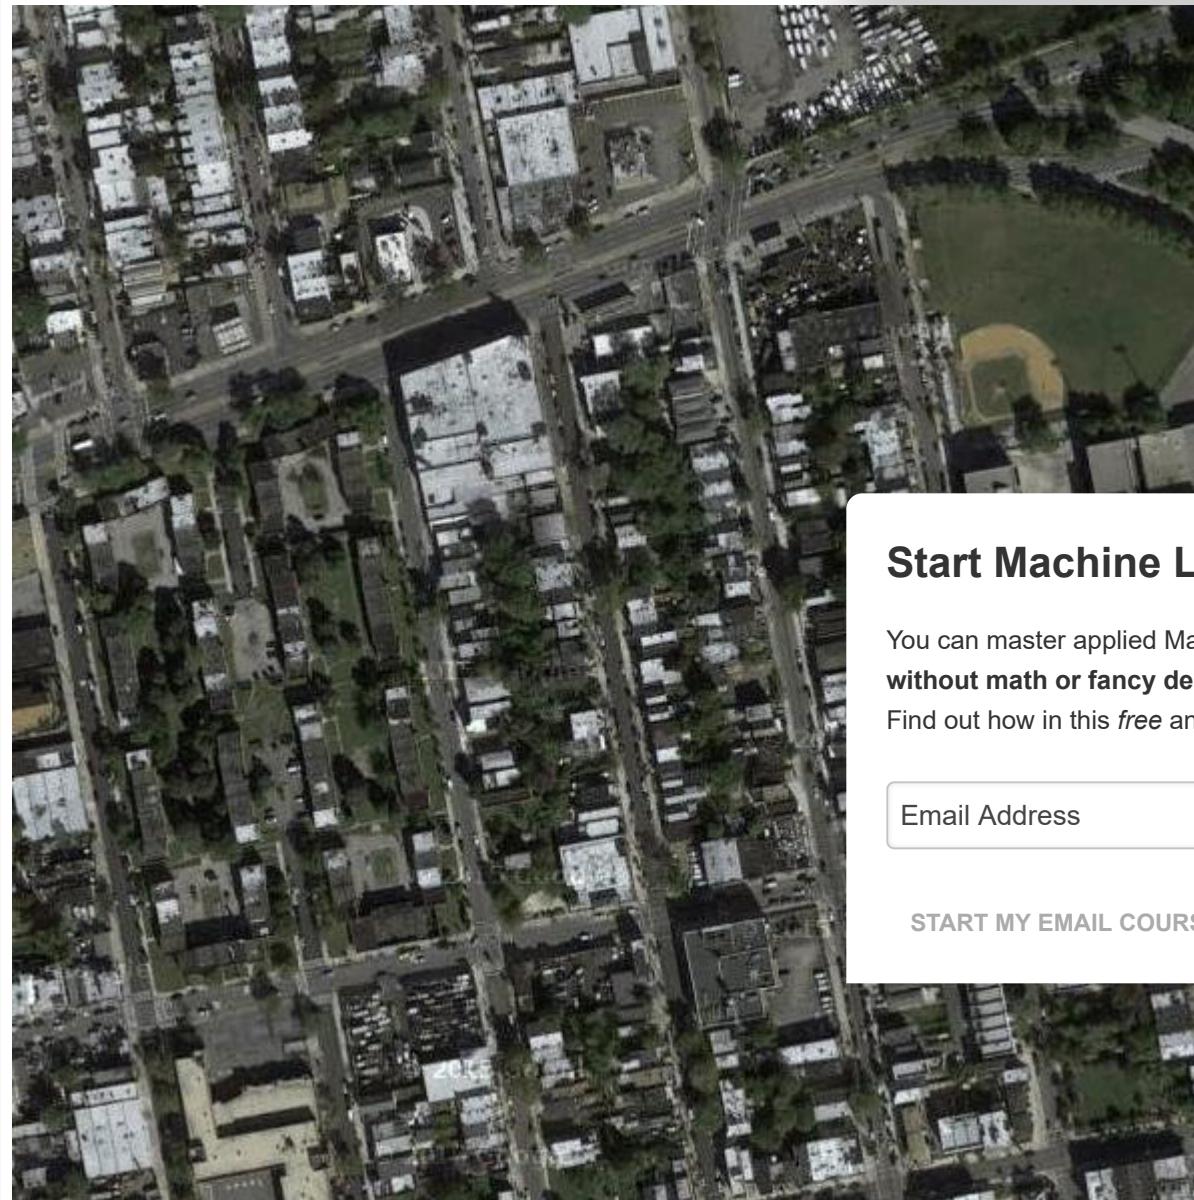
We can select an image from the validation dataset under maps/val and crop the satellite element of the image. This can then be saved and used as input to the model.

In this case, we will use “maps/val/1.jpg”.



We can use an image program to create a rough crop of the satellite element of this image to use as input and save the file as `satellite.jpg` in the current working directory.

Start Machine Learning



Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

Example of a Cropped Satellite Image to Use as Input to the Pix2Pix Model.

We must load the image as a NumPy array of pixels with the size of 256×256, rescale the pixel values to the range [-1,1], and then expand the single image dimensions to represent one input sample.

The `load_image()` function below implements this, returning image pixels that can be provided to the model as input samples.

[Start Machine Learning](#)

```

1 # load an image
2 def load_image(filename, size=(256,256)):
3     # load image with the preferred size
4     pixels = load_img(filename, target_size=size)
5     # convert to numpy array
6     pixels = img_to_array(pixels)
7     # scale from [0,255] to [-1,1]
8     pixels = (pixels - 127.5) / 127.5
9     # reshape to 1 sample
10    pixels = expand_dims(pixels, 0)
11    return pixels

```

We can then load our cropped satellite image.

```

1 ...
2 # load source image
3 src_image = load_image('satellite.jpg')
4 print('Loaded', src_image.shape)

```

As before, we can load our saved Pix2Pix generator model and generate a translation of the image.

```

1 ...
2 # load model
3 model = load_model('model_109600.h5')
4 # generate image from source
5 gen_image = model.predict(src_image)

```

Finally, we can scale the pixel values back to the range [0,1] and plot the result.

```

1 ...
2 # scale from [-1,1] to [0,1]
3 gen_image = (gen_image + 1) / 2.0
4 # plot the image
5 pyplot.imshow(gen_image[0])
6 pyplot.axis('off')
7 pyplot.show()

```

Tying this all together, the complete example of performing an ad hoc image translation with a single image file is listed below.

```

1 # example of loading a pix2pix model and using it for one-off image translation
2 from keras.models import load_model
3 from keras.preprocessing.image import img_to_array
4 from keras.preprocessing.image import load_img
5 from numpy import load
6 from numpy import expand_dims
7 from matplotlib import pyplot

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

```
8
9 # load an image
10 def load_image(filename, size=(256,256)):
11     # load image with the preferred size
12     pixels = load_img(filename, target_size=size)
13     # convert to numpy array
14     pixels = img_to_array(pixels)
15     # scale from [0,255] to [-1,1]
16     pixels = (pixels - 127.5) / 127.5
17     # reshape to 1 sample
18     pixels = expand_dims(pixels, 0)
19     return pixels
20
21 # load source image
22 src_image = load_image('satellite.jpg')
23 print('Loaded', src_image.shape)
24 # load model
25 model = load_model('model_109600.h5')
26 # generate image from source
27 gen_image = model.predict(src_image)
28 # scale from [-1,1] to [0,1]
29 gen_image = (gen_image + 1) / 2.0
30 # plot the image
31 pyplot.imshow(gen_image[0])
32 pyplot.axis('off')
33 pyplot.show()
```

Running the example loads the image from file, creates a translation of it, and plots the result.

The generated image appears to be a reasonable translation of the source image.

The streets do not appear to be straight lines and the detail of the buildings is a bit lacking. Perhaps with further training or choice of a different model, higher-quality images could be generated.

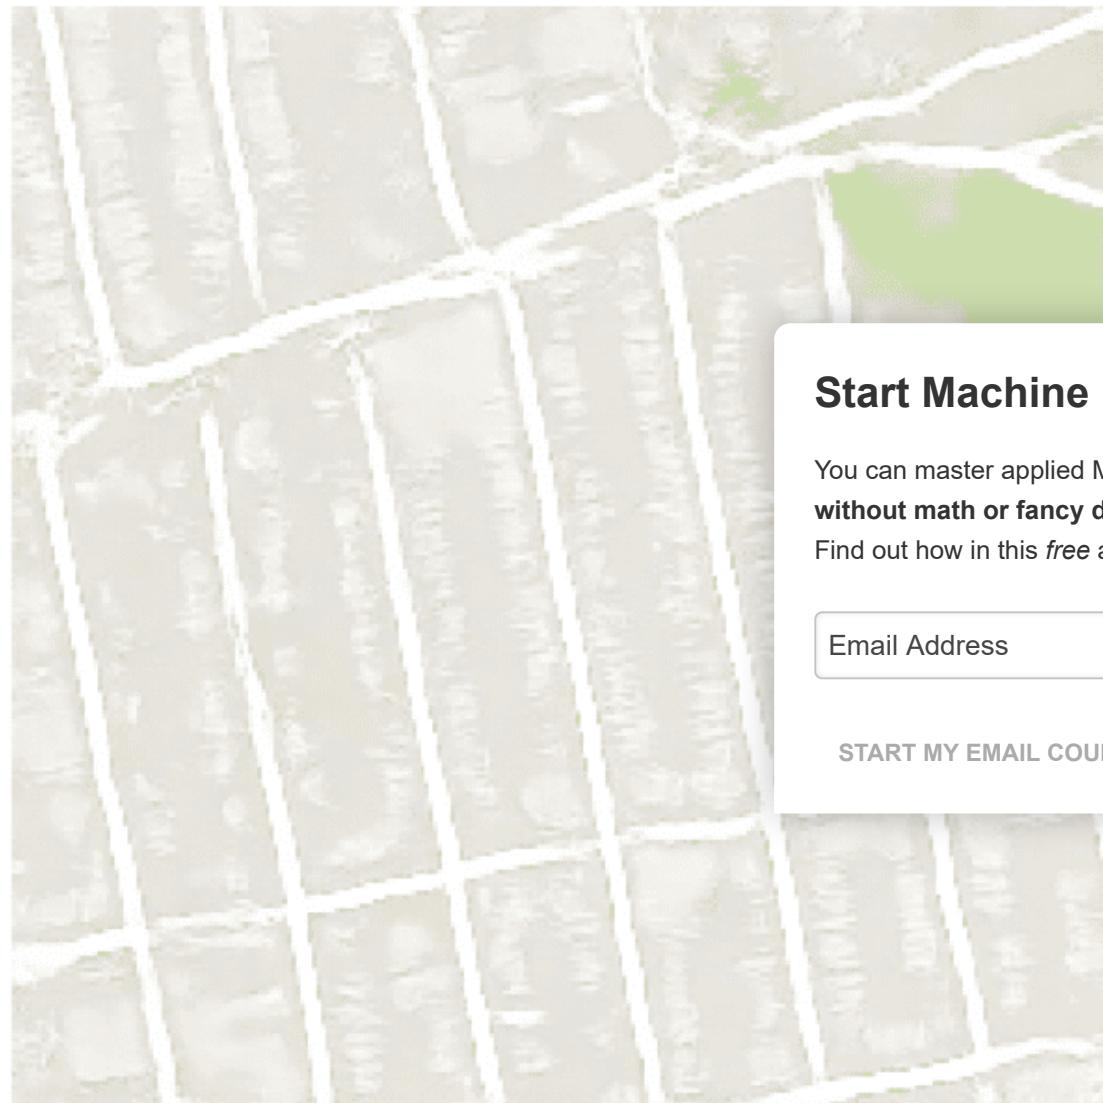
Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning



Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

How to Translate Google Maps to Satellite Images

Now that we are familiar with how to develop and use a Pix2Pix model for translating satellite images to Google maps, we can also explore the reverse.

That is, we can develop a Pix2Pix model to translate Google map images to plausible satellite images. This requires that the model invent or hallucinate plausible buildings, roads, parks, and more.

We can use the same code to train the model with one small difference. We can change the order of the datasets returned from the `load_real_samples()` function; for example:

```

1 # load and prepare training images
2 def load_real_samples(filename):
3     # load compressed arrays
4     data = load(filename)
5     # unpack arrays
6     X1, X2 = data['arr_0'], data['arr_1']
7     # scale from [0,255] to [-1,1]
8     X1 = (X1 - 127.5) / 127.5
9     X2 = (X2 - 127.5) / 127.5
10    # return in reverse order
11    return [X2, X1]
```

Note: the order of X1 and X2 is reversed.

This means that the model will take Google map images as input and learn to generate satellite images.

Run the example as before.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

As before, the loss of the model is reported each training iteration. If loss for the discriminator goes to zero and stays there for a long time, consider re-starting the training run as it is an example of a training failure.

```

1 >1, d1[0.442] d2[0.650] g[49.790]
2 >2, d1[0.317] d2[0.478] g[56.476]
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

```
3 >3, d1[0.376] d2[0.450] g[48.114]
4 >4, d1[0.396] d2[0.406] g[62.903]
5 >5, d1[0.496] d2[0.460] g[40.650]
6 ...
7 >109596, d1[0.311] d2[0.057] g[25.376]
8 >109597, d1[0.028] d2[0.070] g[16.618]
9 >109598, d1[0.007] d2[0.208] g[18.139]
10 >109599, d1[0.358] d2[0.076] g[22.494]
11 >109600, d1[0.279] d2[0.049] g[9.941]
12 >Saved: plot_109600.png and model_109600.h5
```

It is harder to judge the quality of generated satellite images, nevertheless, plausible images are generated after just 10 epochs.

Start Machine Learning ×

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

Email Address

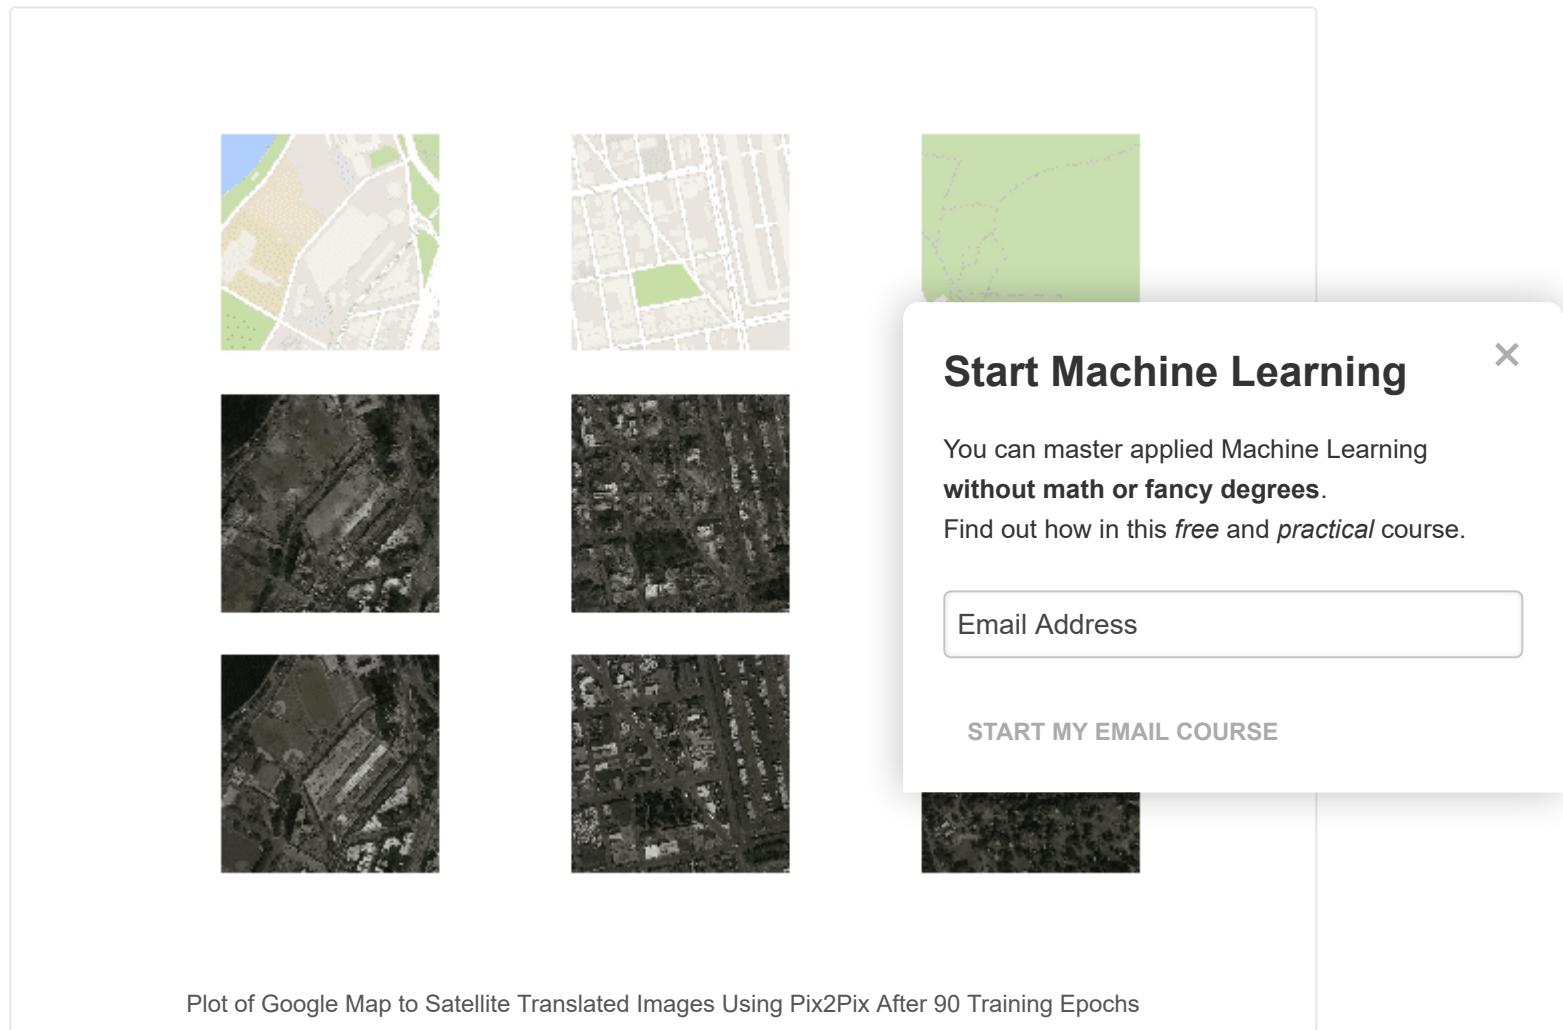
START MY EMAIL COURSE

Plot of Google Map to Satellite Translated Images Using Pix2Pix After 10 Training Epochs

Start Machine Learning

As before, image quality will improve and will continue to vary over the training process. A final model can be chosen based on generated image quality, not total training epochs.

The model appears to have little difficulty in generating reasonable water, parks, roads, and more.



Extensions

This section lists some ideas for extending the tutorial that you may wish to explore.

[Start Machine Learning](#)

- **Standalone Satellite.** Develop an example of translating standalone Google map images to satellite images, as we did for satellite to Google map images.
- **New Image.** Locate a satellite image for an entirely new location and translate it to a Google map and consider the result compared to the actual image in Google maps.
- **More Training.** Continue training the model for another 100 epochs and evaluate whether the additional training results in further improvements in image quality.
- **Image Augmentation.** Use some minor image augmentation during training as described in the Pix2Pix paper and evaluate whether it results in better quality generated images.

If you explore any of these extensions, I'd love to know.

Post your findings in the comments below.

Further Reading

This section provides more resources on the topic if you are looking to go deeper.

Official

- [Image-to-Image Translation with Conditional Adversarial Networks](#), 2016.
- [Image-to-Image Translation with Conditional Adversarial Nets, Homepage](#).
- [Image-to-image translation with conditional adversarial nets, GitHub](#).
- [pytorch-CycleGAN-and-pix2pix, GitHub](#).
- [Interactive Image-to-Image Demo](#), 2017.
- [Pix2Pix Datasets](#)

API

- [Keras Datasets API](#).
- [Keras Sequential Model API](#)
- [Keras Convolutional Layers API](#)
- [How can I “freeze” Keras layers?](#)

Summary

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Start Machine Learning

In this tutorial, you discovered how to develop a Pix2Pix generative adversarial network for image-to-image translation.

Specifically, you learned:

- How to load and prepare the satellite image to Google maps image-to-image translation dataset.
- How to develop a Pix2Pix model for translating satellite photographs to Google map images.
- How to use the final Pix2Pix generator model to translate ad hoc satellite images.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

The image shows the front cover of a book titled "Generative Adversarial Networks with Python". The cover has a blue background. At the top, the title is written in white. Below the title, a subtitle reads "Deep Learning Generative Models for Image Synthesis and Image Translation". At the bottom left, the author's name "Jason Brownlee" is printed. On the bottom right, there is a small icon of a person wearing a VR headset. The overall design is clean and professional.

Develop Generative Adversarial Networks with Python

Develop Your GAN!

...with just a few lines of code

Discover how in my course:
Generative Adversarial Networks with Python

It provides **self-study tutorials** and **end-to-end projects** on:
DCGAN, conditional GANs, image translation, Pix2Pix, CycleGAN
and much more...

Finally Bring GAN Models to your Vision Projects

Skip the Academics. Just Results.

SEE WHAT'S INSIDE

Start Machine Learning

[Tweet](#)[Share](#)[Share](#)

About Jason Brownlee

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

⟨ How to Implement Pix2Pix GAN Models From Scratch With Keras

A Gentle Introduction to CycleGAN for Image Translation ›

258 Responses to *How to Develop a Pix2Pix GAN for Image-to-Image Translation*

Deepanshu Singh August 2, 2019 at 7:33 am #

Amazing tutorial. Detailed and clear explanation of concepts as well as the codes.

Thanks & Regards

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)



Jason Brownlee August 2, 2019 at 2:34 pm #

[REPLY ↗](#)

Thanks!

Sean O'Connor August 3, 2019 at 10:38 am #

[REPLY ↗](#)

[Start Machine Learning](#)

From a digital signal processing viewpoint a weighted sum is an adjustable filter.

Each layer in a conventional artificial neural network has n of those filters and the total compute is a brutal n squared fused multiply accumulates.

A fast Fourier transform is a fixed (nonadjustable) bank of filters, where each filter picks out frequency/phase.

There are other transforms that act as filter banks too, such as the fast Walsh Hadamard transform and these often require far less compute (eg. $n \log(n)$) than a filter bank of weighed sums.

The question then is why not use an efficient transform based filter bank and adjust the nonlinear functions in a neural network by individually parameterizing them?

Ie. change what you adjust:

<https://github.com/S6Regen/Fixed-Filter-Bank-Neural-Networks>

<https://discourse.numenta.org/t/fixed-filter-bank-neural-networks/6392>

<https://discourse.numenta.org/t/distributed-representations-1984-by-hinton/6378/10>



Jason Brownlee August 4, 2019 at 6:21 am #

Perhaps test your alternate approach and compare the results Sean?

Villem Lassmann August 6, 2019 at 7:41 pm #

It seems to me that the discriminator is not a 70×70 PatchGAN, since the 4th layer of the discriminator is a 142×142 PatchGAN. Please correct me if I am mistaken.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee August 7, 2019 at 7:50 am #

REPLY ↗

I believe you are mistaken.

You can learn more about the 70×70 patch gan in greater detail in this post:

<https://machinelearningmastery.com/how-to-implement-pix2pix-gan-models-from-scratch-with-keras/>

Start Machine Learning

REPLY ↩

Villem Lassmann August 7, 2019 at 8:17 pm #

That example has the same structure, 6 layers of Conv2D (including the last one). But when looking at the beginning of the post where You are calculating the receptive field with 5 layers of Conv layers. The calculation also states that there are only 3 layers of Conv2D with a stride of 2. I believe that the layer named C512 should be the second to last layer.

**Jason Brownlee** August 8, 2019 at 6:40 am #

REPLY ↩

I believe the implementation matches the official implementation described here:

<https://github.com/phillipi/pix2pix/blob/master/models.lua#L180>

Vasudevakrishna August 20, 2019 at 3:50 pm #

Thanks for the tutorial.

My question is in original paper they are giving the direction as configurable parameter.

But in your implementation I am unable to see that one.

How can do that for both direction.

Please explain

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

X

**Jason Brownlee** August 21, 2019 at 6:34 am #

REPLY ↩

I show how to translate images in both directions in the above tutorial.

Alex August 29, 2019 at 11:52 pm #

REPLY ↩

Many thanks for this amazing tutorial!

PS “There are 1,097 images”... and then there are saves every 10970 steps, and 109700 ste

Start Machine Learning



Jason Brownlee August 30, 2019 at 6:24 am #

REPLY ↗

Thanks.

Fixed.

Salman Sajd September 1, 2019 at 7:13 am #

REPLY ↗

Thanks for an amazing tutorial

How we use GAN for motion transfer or which type of GAN will best for Motion Transfer?



Jason Brownlee September 2, 2019 at 5:24 am #

I don't know off hand sorry, perhaps try a search on scholar.google.com

Lin September 6, 2019 at 9:52 pm #

Hello, thanks for the great article.

I have one question, but why you scale the image to [-1, 1] instead of [0, 1]?

Does this make the model behave differently?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee September 7, 2019 at 5:29 am #

REPLY ↗

Because the generator generates pixels in that range, and the discriminator must "see" all images pixels in the same range.

Start Machine Learning

The choice of -1,1 for pixels is a gan hack:

<https://machinelearningmastery.com/how-to-code-generative-adversarial-network-hacks/>

Dang Tuan Hoang September 12, 2019 at 1:45 pm #

REPLY ↗

Hi sir, is it possible to train this model with inputs and output of different sizes?

For example, I have 3 image a,b,c with size 50x50x3. I want the model to generate c from a,b. First I append a and b to get d with size 50x100x3, then use d as input, c as output



Jason Brownlee September 12, 2019 at 1:49 pm #

Yes, you can use different sized input and output, although the u-net will

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee September 17, 2019 at 6:24 am #

Sorry, I don't have the capacity to prepare custom code for you.

Perhaps experiment with adding/subtracting groups of layers to the decoder part of the model and see the effect on the image size?

Dang Tuan Hoang September 17, 2019 at 12:37 pm #

Start Machine Learning

I know you are very busy so I didn't ask for custom code, I just need something to start with. Thank for the suggestion sir !



Jason Brownlee September 17, 2019 at 2:34 pm #

Perhaps start with just the function that defines the model and try playing around with it.

David September 27, 2019 at 9:17 am #

REPLY ↗

Hi Jason,

First of all, thank you very much for posting this tutorial, I learned a lot from it.

I have a question.

Do u think if I leave the picture resolution as it is rather than compressing them.

The performance is gonna be better? As my pictures between translation is very minor.

Thank you!

David

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

REPLY ↗



Jason Brownlee September 27, 2019 at 1:15 pm #

Thanks, I'm happy that it helped.

Interesting idea. You mean likely working with TIFF images or other loess-less formats?

Probably not, but perhaps test it to confirm.

Alice October 16, 2019 at 6:16 pm #

REPLY ↗

Hello,

Start Machine Learning

How can I increase speed of training? It uses very small portion of gpu memory.



Jason Brownlee October 17, 2019 at 6:26 am #

REPLY ↗

Some ideas:

Use less data.

Use a smaller model.

Use a faster machine.

Alice October 17, 2019 at 8:33 am #

I am using a machine with 8 gpus (8 X p4000) 😊

I mean, for example, while training on darknet, changing batch size directly affects gpus each gpu. And batch size doesn't affect it. So I need an adjustment just like on darknet. Thanks



Jason Brownlee October 17, 2019 at 1:50 pm #

I see, I'm not sure I can help sorry.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Samuel October 17, 2019 at 8:30 am #

REPLY ↗

Swapping out the training data for the SEN1-2 dataset had amazing results. I can now translate Sentinel 1 images to RGB Sentinel 2! Many thanks for such a thorough tutorial.

Start Machine Learning



Jason Brownlee October 17, 2019 at 1:49 pm #

REPLY ↗

Well done!

I would love to see an example of a translated image.

Mohammad October 20, 2019 at 12:03 pm #

REPLY ↗

Hi Jason,

Thank you so much for your great website, it is fantastic.

I was wondering what your opinion is about the future research direction for this area of research.

Thanks



Jason Brownlee October 21, 2019 at 6:14 am #

X

Thanks.

Sorry, I don't have thoughts on research directions – I try to stay focused on the industrial

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Michael October 28, 2019 at 9:08 am #

REPLY ↗

Awesome tutorial on Pix2Pix. Your other GAN articles were great too and very helpful. After reading your tutorials, I was able to implement my own Pix2Pix project. All the code is on my GitHub. <https://github.com/michaelnation26/pix2pix-edges-with-color>



Jason Brownlee October 28, 2019 at 1:18 pm #

REPLY ↗

Thanks.

Start Machine Learning

Well done, that is very impressive!

sss November 17, 2019 at 10:43 pm #

REPLY ↗

- python version: 3.6.7
- tensorflow-gpu version: 2.0.0
- keras version: 2.3.1
- cuDNN version: 10.0
- CUDA version: 10.0

mnist_mlp.py (https://github.com/keras-team/keras/blob/master/examples/mnist_mlp.py) works perfectly but code which is given below gives me this error:

```
tensorflow.python.framework.errors_impl.UnknownError: Failed to get convolutional kernel from CUBLAS handle. Could not find an available CUDA stream. This often means cuDNN failed to initialize, so try looking to see if a warning log message like the following was present in your terminal:
[[node conv2d_7/convolution (defined at C:\Users\ACSECKIN\Anaconda3\envs\tensorflow\lib\site-packages\tensorflow_core\python\framework\ops.py:1751) ]] [Op:__inference_k
```

Function call stack:
keras_scratch_graph

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

X



Jason Brownlee November 18, 2019 at 6:45 am #

REPLY ↗

Perhaps try running on the CPU first to confirm your environment is working?

sss November 18, 2019 at 6:04 pm #

REPLY ↗

Code works with CPU. But I want to run in the GPU to complete in less time. CPU time is about 16 hours. I am open to any alternative to decrease the training time.

Start Machine Learning



Jason Brownlee November 19, 2019 at 7:38 am #

REPLY ↗

That is odd. I ran the examples on GPU without incident, specifically on EC2.

Perhaps try that?

Jack November 27, 2019 at 6:12 pm #

Based on my experience, I receive this error message when my GPU does not have enough memory to handle the process. Maybe try reducing the computational workload by using a smaller image? If not, use CPU if you are fine with it.



Jason Brownlee November 28, 2019 at 6:33 am #

Great suggestions!

Ivan November 23, 2019 at 9:05 pm #

Hi guys,

It takes 8 hours to train the model on GPU (Floydhub).

But several different models were saved in the process.

Can you explain why?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee November 24, 2019 at 9:17 am #

REPLY ↗

Perhaps try training on a large ec2 instance, it is much faster.

Models are saved periodically during training, we cannot know what a good model will be

Start Machine Learning

Ivan November 25, 2019 at 6:52 pm #

REPLY ↗

Thanks for the tip!

Thought that more epochs give better results, since GAN's cannot converge... so by theory no over fitting.. but I am new in the field, so I will look more into it 😊



Jason Brownlee November 26, 2019 at 6:00 am #

REPLY ↗

Not with GANs. Perhaps start here:

<https://machinelearningmastery.com/start-here/#gans>

Jack November 26, 2019 at 2:07 pm #

Thank you for your awesome post, it is really detailed and helpful! I have a question: images are single channel and the maximum and minimum pixel values vary for each image, should I go about normalizing the images? Should I take the maximum of the whole batch of samples for each sample and normalize each individually?

Also, when translating new images, what would be the values of image? Would it be between -1 to 1? If yes, how should I "denormalize" the values to the original? Thank you for your help!

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee November 26, 2019 at 4:10 pm #

REPLY ↗

Thanks.

I recommend selecting a min and max that you expect to see for all time and use that to scale the data. If not known or it cannot be known, use the limits of the domain (e.g. 0 and 255).

Start Machine Learning



Jack November 27, 2019 at 6:11 pm #

[REPLY ↗](#)

Thank you for your suggestions! How would you suggest I “de-normalize” the data during testing? Should I use the same range (I am taking the range from the training data) and reverse the process on the test data?



Jason Brownlee November 28, 2019 at 6:32 am #

[REPLY ↗](#)

Yes.



Syd Rawat December 1, 2019 at 2:50 pm #

Hi Jason,

Thank you very much for sharing such an in-depth analysis of Pix2Pix GAN. It is really helpful for my CS background. I thought of applying this for solving and inverse problems in Digital Holography. I have got some preliminary results I have got. As you know, the output of the model is a translated image, here I am looking for an image quality metric such as SSIM. Do you have any suggestions?

Thank You,

PS: As this post helped me enormously, I would like to cite your works on GANs in the future.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

Jason Brownlee December 2, 2019 at 5:58 am #

[REPLY ↗](#)

You're welcome.

That sounds very cool! Perhaps one of the metrics here would be helpful:

<https://machinelearningmastery.com/how-to-evaluate-generative-adversarial-networks/>

Start Machine Learning

[REPLY ↗](#)**Anthony** December 15, 2019 at 1:24 am #

Hi! Currently implementing this with images with shape (256, 512, 3) and keep running to an error as follows:

"ValueError: A target array with shape (1, 16, 16, 1) was passed for an output of shape (None, 16, 32, 1) while using as loss binary_crossentropy. This loss expects targets to have the same shape as the output."

I assume that this is due to the downsampling? Any help would be appreciated

[REPLY ↗](#) **Jason Brownlee** December 15, 2019 at 6:08 am #

Perhaps start with square images, get it working, then try changing to rectangular.

[REPLY ↗](#)**Anthony Mäkelä** December 15, 2019 at 6:21 am #

Hmm, alright! Could you explain why you use target_size=(256, 512) instead

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)**Jason Brownlee** December 16, 2019 at 6:03 am #

The images are 256×512 – as they contain the input and output images together.

We load them and split them into two separate square images 256×256 when working with the GAN.

[REPLY ↗](#)**MK** December 16, 2019 at 11:45 pm #

The discriminator error seems to be going to zero pretty quick, any tips to avoid this?

[Start Machine Learning](#)



Jason Brownlee December 17, 2019 at 6:36 am #

REPLY ↗

Perhaps try running the example a few times and continue running – it may recover.

Avdudia December 17, 2019 at 4:58 am #

REPLY ↗

Thank you for this tutorial and simple code. I used it to perform image-to-image translation from Köppen–Geiger climate classification map (https://en.wikipedia.org/wiki/K%C3%B6ppen_climate_classification) to real satellite data, with truly amazing results, but I have a question.

In my strategy I create near one thousand pairs of 256×256 tiles from the Köppen–Geiger map (present in the Wikipedia article above), and a high-resolution satellite map of the Earth. In order to minimize deformation on tiles pairs near poles, I create pairs of image for GAN training (see <https://photos.app.goo.gl/eGvpXghUtCB9kqkX6>).

I trained the GAN until the end (n_epochs=100) with amazing results. Using training data given (<https://photos.app.goo.gl/a4EV6Gh15hAnYokm7>). Even with hand-painted or with source image in Geiger colormap, results are very nice (<https://photos.app.goo.gl/eGbFmTH7YqYi4xfu5>).

However I noticed that the result lacked of “relief” effect. Moreover, on large landmasses where the heightmap noticeably affects the satellite view (e.g Tibetan Plateau or the Grand Canyon), the model results are not good.

As the climate map is composed of only 29 different indexed colors (plus the one I added for clouds), I decided to store the heightmap on the first channel of the input image, and the climate index on the second channel. The third channel is kept unused. It results in a Red-Green image where the Red channel is the heightmap and the Green channel is normalized climate index (see <https://photos.app.goo.gl/cN1cmCNLSXwwqzNB9>).

The problem is that training with this input images give bad results compared to my first try (only climate date). Results were already convincing after 30 epochs in my first try, with smooth transition between climates, why here the boundaries are clearly visible in generated images (see <https://photos.app.goo.gl/Q1vjjeY8ewWrCZYv5>).

I tried to run the training several times to ensure that it was not purely bad luck, with the same result.

I don't understand because climate index can clearly be stored on one channel without information loss, and heightmap provides additional data, so it should improve the results. Is it simply because it needs more epochs ?

Thank you in advance and sorry for the long post and for my english, it is not my native language^^

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee December 17, 2019 at 6:42 am #

REPLY ↗

Well done!

Very cool application.

Two thoughts off the cuff. One would be to make image synthesis conditional on two input images (source and the height map). A second would be to have 2 steps – one step to generate the initial translation and a second to improve the generated image with a height map.

I'm eager to hear how you go!

Avdudia December 18, 2019 at 7:32 am #

Thank you very much ! The aim is to develop a tool for worlbuilding and creating Atrifexian's tutorials <https://www.youtube.com/watch?v=5ICbxMZJ4zA&t=1s>).

I used the idea of using R and G channels for heightmap and climate following this thread <https://github.com/junyanz/pytorch-CycleGAN-and-pix2pix/issues/498>. They recommend your code is limited to 3 channels and as I'm a complete beginner I still don't know how

However it seems indeed that training on more epochs actually gives good results without restarting it with a limit of 1000 epochs. However I have to redo the first 100, as I run the unstable (I only managed to reach 100 epochs twice).

Do you have a tutorial on how to make complete checkpoints in order to continue the training in case of crash ? If I understand well, your summarize_performance function only save the generator model, so we should have to save the whole gan_model and reload it for later training. Do you have documentation or examples concerning this ?

Thank you so much for your tutorial. I'll keep you informed on later developments !



Jason Brownlee December 18, 2019 at 1:28 pm #

REPLY ↗

Yes, the above code already saves the model every n steps.

See the summarize_performance() function.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

JS December 18, 2019 at 2:06 am #

REPLY ↗

Should i approach this the same way if i have images containing white backgrounds? Similar to the edges2shoes dataset?



Jason Brownlee December 18, 2019 at 6:10 am #

REPLY ↗

Perhaps test it with a prototype?

JS December 18, 2019 at 8:01 pm #

For some reason i end up only with blank white images...

Sirine December 25, 2019 at 5:56 am #

Hello,

Thanks for the wonderful tutorial, Please how can I adapt the generator and the descriminator matrix(2,64)

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee December 25, 2019 at 10:37 am #

REPLY ↗

Sorry, I don't follow.

Elbek Khoshimjonov January 2, 2020 at 7:59 pm #

Start Machine Learning

Thanks Jason for great post!

I have tried this code, but images do not appear to be good enough, and discriminator loss becomes 0 after 10-15 epochs.



Jason Brownlee January 3, 2020 at 7:27 am #

REPLY ↗

Perhaps try using a different final model or re-fit the model?

Arsal January 28, 2020 at 10:55 pm #

REPLY ↗

I want to continue training from the last checkpoint stored. Can you help me in resum



Jason Brownlee January 29, 2020 at 6:36 am #

Yes, load the model as per normal and call the train() function.

Arsal January 28, 2020 at 10:57 pm #

I have a dataset consisting of 216 images. I trained for 100 epochs but unfortunately the results are not good. Can you help me how can I improve the results?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee January 29, 2020 at 6:36 am #

REPLY ↗

Yes, try some of these suggestions:

<https://machinelearningmastery.com/how-to-code-generative-adversarial-network-hacks/>

Start Machine Learning

REPLY ↗

Laura January 30, 2020 at 3:19 am #

Thank you for this wonderful tutorial! It has been extremely helpful. I was wondering if you had considered data augmentation?

**Jason Brownlee** January 30, 2020 at 6:56 am #

REPLY ↗

For GANs, not really, in general, yes:

<https://machinelearningmastery.com/how-to-configure-image-data-augmentation-when-training-deep-learning-neural-networks/>

Laura January 30, 2020 at 3:43 am #

Thanks!

This is probably a newbie question, but I am new to GANs. In my limited experience with deep learning, I would use validation sets during the training process to sort of evaluate how well it was “learning” > I then had another dataset I called the “test set” which I used to evaluate the final model after training was complete. Here it seems like you don’t use any validation during the training process. And then you use the test set as the validation dataset. Is that something unique to GANs or can validation be included in the training process?

**Jason Brownlee** January 30, 2020 at 6:58 am #

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

No, I often use tests sets for validation to make tutorials simpler:

<https://machinelearningmastery.com/faq/single-faq/why-do-you-use-the-test-dataset-as-the-validation-dataset>

Michael February 1, 2020 at 12:34 am #

REPLY ↗

Hi! Im running this on Titan V and it seems to be running extremely slow. Any ideas as to why?

Start Machine Learning



Jason Brownlee February 1, 2020 at 5:56 am #

REPLY ↗

Training GANs is slow...

Perhaps check you have enough RAM and GPU RAM?

Perhaps compare results on an ec2 with a lot of RAM?

Perhaps adjust model config to use less RAM?

Lisa October 28, 2020 at 8:57 pm #

REPLY ↗

I suggest using google colab. You can use their GPUs for training. It's much faster.



Jason Brownlee October 29, 2020 at 8:02 am #

GPUs are practically a requirement when working with GANs.

pramod kumar February 9, 2020 at 2:42 am #

sir can i know how you downloaded the images of satellite and maps can you please help me to download my own dataset for this project



Jason Brownlee February 9, 2020 at 6:24 am #

REPLY ↗

See the section of the above tutorial “Satellite to Map Image Translation Dataset” on how to download the dataset.

Vaibhav Vijay kotwal February 24, 2020 at 6:28 pm #

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Start Machine Learning

Hi Jason,

From the theory, we understand that discriminator learns the objective loss function. However referring to define_GAN(), line 56 in the code, I am not able to see the object loss learnt by discriminator getting passed to GAN model. I see that the model does not converge as expected

Thanks and Regards,

Vaibhav Kotwal



Jason Brownlee February 25, 2020 at 7:42 am #

REPLY ↗

We create a composite model that combines G and D so that G is trained via inverse corrections to D.

Perhaps see this:

<https://machinelearningmastery.com/how-to-code-the-generative-adversarial-network-training/>

Firat Erdem February 26, 2020 at 10:49 pm #

Thanks for the great tutorial. I need help with something. I want to see accuracy metric for the training process. And I want to see this for each epoch, not for each steps. Like a standard things to code ? I couldn't apply it because it is different than standard CNN codes. I would re

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee February 27, 2020 at 5:50 am #

REPLY ↗

Accuracy is a bad metric for GANs.

See this:

<https://machinelearningmastery.com/how-to-evaluate-generative-adversarial-networks/>

Tom March 8, 2020 at 5:13 am #

REPLY ↗

Can pix2pix Gan save and load again without training again

<https://machinelearningmastery.com/how-to-develop-a-pix2pix-gan-for-image-to-image-translation/>

Start Machine Learning



Jason Brownlee March 8, 2020 at 6:16 am #

REPLY ↗

Yes, see this tutorial for an example:

<https://machinelearningmastery.com/how-to-develop-a-pix2pix-gan-for-image-to-image-translation/>

Tom March 8, 2020 at 6:37 am #

REPLY ↗

Yes, thank for your reply!! And I want to ask another questions!! Can the other kind of GAN versions save and load?



Jason Brownlee March 9, 2020 at 7:08 am #

Yes, they are all Keras models that can be saved and loaded.

Ehsan Karimi March 12, 2020 at 2:10 am #

Hi Jason,

Thanks for the great tutorial. I have a problem with image scales. In first step, after splitting the input images, I check the image size, instead of 256*256 pixel they are 134*139 with background. Also, at translation a given standalone image using by model step, the output should be 256*256 same as input, but I get 252*262 output again with background.

I was wondering if you would mind letting me know where is the problem?

Thanks in advance

Ehsan

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee March 12, 2020 at 8:51 am #

REPLY ↗

Start Machine Learning

I don't know the cause of your fault. Sorry.

Paolo March 13, 2020 at 8:13 am #

REPLY ↗

Great work Jason. Just one question: do you believe that this approach could work using a RGB satellite image against its mask image, to make some kind of image segmentation ?

Thanks in advance



Jason Brownlee March 13, 2020 at 8:24 am #

Perhaps try it. Prototypes are a fast way to get answers.

Paolo March 13, 2020 at 8:31 am #

I mean, the mask image would have just 2 colors (yes/not) ... this was my co

Iqbal March 19, 2020 at 9:05 pm #

It is an interesting article publish here. I am new using this, i want some question for the first script for clear explanation :

1. I saw the loaded data is maps in train and test folder. I want to know which 3 sample was loaded from the folder train? because the results was : (1096, 256, 256, 3) (1096, 256, 256, 3). I understand 1096 is the certain amount of image in that folder. and 256 I still dont understand because when I open picture 256 is not the same as the it was loaded.

2. I saw the folder contain image in train and test. I want to ask the train, example 1.jpg it contains two image from satellite. May I know how to develop the left picture and the right picture or it develop itself? Also in test does it develop itself or have to save it first?

Need some explanation for preparing using it in the future. Thank you

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee March 20, 2020 at 8:43 am #

REPLY ↗

The images are 256×256 squares with 3 color channels. You can learn more about loading images here:
<https://machinelearningmastery.com/how-to-load-and-manipulate-images-for-deep-learning-in-python-with-pil-pillow/>

Sorry, I don't follow your second question, perhaps you can elaborate?

Iqbal March 20, 2020 at 9:51 am #

REPLY ↗

thank you for replying. For my second question I saw the folder train and test contains an image. So the question come up :

- 1.Does the image build itself?
2. If No, how do you make the images side by side that contains two image in one image?

Thank you very much.



Jason Brownlee March 20, 2020 at 1:15 pm #

The tutorial shows how to load the images and prepare them for modeling.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

REPLY ↗

yasser March 22, 2020 at 10:12 am #

hi thank you for your work !

I need your help,I need the same model but the input of the generator is one channel and not three .
I have tried to change it but it does'nt work .thank you



Jason Brownlee March 23, 2020 at 6:09 am #

REPLY ↗

Sorry to hear that.

Start Machine Learning

Perhaps confirm that your images are grayscale (1 channel), then change the model to expect 1 channel via the input shape.

runnergirl March 23, 2020 at 5:53 am #

REPLY ↗

Hi! I wanted to train myself. I prepared them just like in this tutorial. Size X1 and X2 are the same. Data display works. But I get this error:

'Got inputs shapes: %s' % (input_shape))

ValueError: A Concatenate layer requires inputs with matching shapes except for the concat axis. Got inputs shapes: [(None, 2, 2, 512), (None, 1, 1, 512)]

What have I done wrong?



Jason Brownlee March 23, 2020 at 6:16 am #

I don't know sorry.

Mick October 28, 2020 at 10:26 pm #

Have you been able to solve this problem. I get the same issue.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

GoComputing October 30, 2020 at 10:36 pm #

REPLY ↗

This is probably because your input size is not divisible by 256.

ouis March 24, 2020 at 1:04 am #

REPLY ↗

Start Machine Learning

hello I have this problem I dont know why :

ValueError: Graph disconnected: cannot obtain value for tensor Tensor("input_14:0", shape=(?, 256, 256, 3), dtype=float32) at layer "input_14". The following previous layers were accessed without issue: ['input_15']



Jason Brownlee March 24, 2020 at 6:05 am #

REPLY ↗

Perhaps confirm that your keras and tensorflow versions are updated?

Harshit April 19, 2020 at 2:37 am #

Hi. I tried to use a different dataset using this code. Specifically the edges2shoes dataset. Everytime i ran into memory error. My ram is 16GB still that was not enough. I managed to create a smaller version of the dataset. Also could you be kind enough to make tutorial of Tensorflow/Keras of pix2pixHD since it is more complex than normal pix2pix.



Jason Brownlee April 19, 2020 at 6:01 am #

REPLY ↗

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Perhaps use a sample of the data?

Perhaps use progressive loading?

Perhaps use an ec2 instance wth more ram?

Thanks for the suggestion.

Harshit April 19, 2020 at 11:48 am #

REPLY ↗

Thanks I think your tutorial here <https://machinelearningmastery.com/how-to-load-large-datasets-from-directories-for-deep-learning-with-keras/> will be helpful in my case.

Start Machine Learning

I am using my own machine not AWS instance.



Jason Brownlee April 19, 2020 at 1:18 pm #

REPLY ↗

Great!

Phil April 30, 2020 at 11:23 pm #

REPLY ↗

Hi Jason. Great Article. Good explanation. Your articles gave me a good overview and starting point when I started developing my own networks. But I have two questions:

- 1.) From what I can see the original code on Github seems to be slightly different to your code. In your code you pass data from an encoder layer to a decoder layer. On Github data from an encoder layer passed to a decoder layer. In your code the data is passed directly after the convolution(or batch norm/dropout), in contrary to the solution on Github.
 - 2.) What does the flag 'training=True' do when calling batch normalization layer or dropout layer?
- Thanks in advance.

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



Jason Brownlee May 1, 2020 at 6:40 am #

Perhaps. I thought I had the architecture spot on based on the paper and the released code.

Training=True causes the layer to think it is always in training model. E.g. normally batchnorm and dropout operate differently in training vs inference. More here:

<https://machinelearningmastery.com/how-to-implement-pix2pix-gan-models-from-scratch-with-keras/>

ElenaRR May 3, 2020 at 4:08 am #

REPLY ↗

Hi Jason,
thank you very much for this tutorial, it's awesome!

Start Machine Learning

I have the problem that you mentioned at the end of the article. D1 loss goes to zero after 80-90 steps. Could you explain me why this happens and how can I solve it?

In addition to this, I can see that only one image is used in every iteration (one real, one fake) where n_batch = 1. Shouldn't we use more than one pair of images to train in each step?

```
# select a batch of real samples
[X_realA, X_realB], y_real = generate_real_samples(dataset, n_batch, n_patch)
```

Thank you very much!



Jason Brownlee May 3, 2020 at 6:18 am #

Yes, this is probably a failure mode:

<https://machinelearningmastery.com/practical-guide-to-gan-failure-modes/>

You can try fitting the same model again, save models to file along the way.

You can try tuning the model architecture or training algorithm:

<https://machinelearningmastery.com/how-to-code-generative-adversarial-network-hacks/>

REPLY ↗

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

X

ElenaRR May 3, 2020 at 7:23 pm #

So, the model as it is in this example is not going to work properly? This link is not a pixtopix architecture. I tried also the example and work perfectly for cGAN and GAN with fmnist, but the problem is this Pixtopix architecture.

REPLY ↗



Jason Brownlee May 4, 2020 at 6:19 am #

The example does work correctly, but you may need to train it a few times.

GANs are unstable by definition.

Start Machine Learning

 **ElenaRR** May 3, 2020 at 7:39 pm #

REPLY ↗

I've seen that there are specific codes in your book about this. Is a more complete example for pixtopix? Or is it the same?

**Jason Brownlee** May 4, 2020 at 6:19 am #

REPLY ↗

The examples of pix2pix in the book are based on this example.

 **shami** May 11, 2020 at 1:02 am #

long live 100 years superb tutorial with clear explanation

**Jason Brownlee** May 11, 2020 at 6:02 am #

Thanks!

 **bobbyP** May 14, 2020 at 2:29 pm #

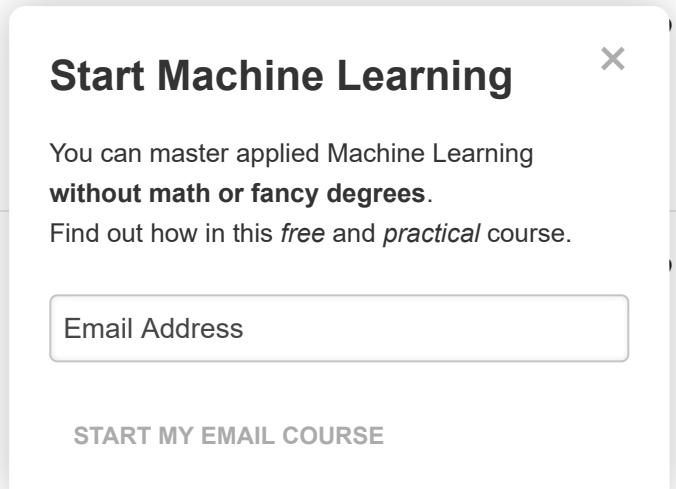
anyway to use a keras data loader/generator for on the fly image loading during training? say for example your training size is really large and loading all at once would result in out of memory errors? thanks so much for all your tutorials, they are incredible!

**Jason Brownlee** May 15, 2020 at 5:55 am #

REPLY ↗

Yes, see this:

[https://machinelearningmastery.com/how-to-load-large-datasets-from-directories-for-deep learning/](https://machinelearningmastery.com/how-to-load-large-datasets-from-directories-for-deep-learning/)

 Start Machine Learning

cans May 20, 2020 at 6:12 am #

REPLY ↗

Hi Sir,

First of all i read your all tutorials. You are helping me more than my consultant. Thank you so much.

I am new at Gans. Sorry for my qusetions. But i cant understand how i test this model?

I will use validation set okey.

After training model i wont give target, just give source image?I cant get it.

Or only i load model and train with validity set?

Omg i cant explain myself. I hope u are understand me.



Jason Brownlee May 20, 2020 at 6:29 am #

You're welcome!

Evaluating GANs is challenging. We do not use a validation set. Instead we generate images enough.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

cans May 21, 2020 at 2:39 am #

thank you for your answer Sir,

i wanna try pix2pix gans for image enhancement.

I'll use source images are low contrast, targets are high contrast ,what do you think? i'm trying to improve thermal image.

I hope this system will work.



Jason Brownlee May 21, 2020 at 6:21 am #

REPLY ↗

Sounds great!

Start Machine Learning

Gruhit Patel May 25, 2020 at 1:44 pm #

REPLY ↗

Here in Discriminator what is the need for Concatenating Source and Target images ? What effect would it have ?



Jason Brownlee May 26, 2020 at 6:12 am #

REPLY ↗

Sorry, which section are you referring to? Where (which section/line?) do we concatenate images?

Gruhit Patel May 27, 2020 at 2:48 pm #

In discriminator. where we are concatenating Source Image and target Image

Actually I'm building a GAN model color transformation from Gray to RGB.

My discriminator and Generator model's loss falls to zero. So wanted to know that why
discriminator. And if you have any advice for my model than tell it too.

THANKS in advance...

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee May 28, 2020 at 6:09 am #

Here we are training a conditional model, e.g. generate a target image conditional on the source image.

E.g. it is the purpose of the model.

Ujjayant Sinha May 27, 2020 at 12:54 am #

REPLY ↗

Hello. I compared the images from the summarize_performance() to predictions on unseen ones, which turned out to be quite horrible.
Can you suggest some ways to tackle this problem ?

Start Machine Learning



Jason Brownlee May 27, 2020 at 7:58 am #

REPLY ↗

Try training the model a few times, save many times during each run, choose a model that generates good images.

Zsolt Lipcsei May 27, 2020 at 8:50 am #

REPLY ↗

Hi,

How can I use the generator model to predict any size of images? I mean not just squares sizes. Is there any way at all?



Jason Brownlee May 27, 2020 at 1:26 pm #

You will have to change the generator/discriminator and also the training dataset

Muhammad Ammar Malik June 12, 2020 at 2:49 am #

Thank you for the awesome post. I have 2 questions if you can answer please.

First question, you have mentioned:

"In this case, we will use the model saved at the end of the run, e.g. after 10 epochs or 109,600 training iterations."

Shouldn't the training iterations be 10,960 after 10 epochs.

Second question, what is the rationale behind using random index to generate real and fake samples? Why can't we simply iterate over all the samples one by one to make sure no image is missed or used more than once in 1 training step?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee June 12, 2020 at 6:18 am #

REPLY ↗

Images are generated after every 10 epochs, it runs for 100 epochs, meaning we used the model after 100 epochs. Fixed.

Start Machine Learning

We can do it for all image, I wanted to work with one image, to show we can use the model ad hoc. Readers often find that step confusing so I must demonstrate it.

Riccardo June 18, 2020 at 11:27 pm #

REPLY ↗

Hi!

thank you for the great tutorial, you helped me a lot!

i have just a question: am i doing something wrong or is it normal that for a X input i do not have a unique Y output.

Let me explain better: if i repeat n times the prediction i get n different Y images (i'm checking pixels differences).

I'm translating this example to another application and having the exact same output everytime will make it works.

I tried to look for a random noise vector or something like that but it seems that this is not the



Jason Brownlee June 19, 2020 at 6:14 am #

Recall we are using a GAN, so we have two models, the first predicts whether in images conditional on another image.

If you are new to GANs perhaps start here:

<https://machinelearningmastery.com/what-are-generative-adversarial-networks-gans/>

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Riccardo June 18, 2020 at 11:29 pm #

REPLY ↗

EDIT:

" I'm translating this example to another application and having the exact same output will make it works" *

Riccardo June 18, 2020 at 11:31 pm #

REPLY ↗

Start Machine Learning

Hi!

this tutorial helped me a lot!

i have just a question: is it normal that if i repeat n times the prediction, with the same input, i have n different outputs?

I'm checking directly changes in pixel values in Y outputs.



Jason Brownlee June 19, 2020 at 6:15 am #

REPLY ↗

Yes, this is expected given the stochastic nature of some of the layers.

Riccardo June 19, 2020 at 5:35 pm #

thank you for the reply!

I'm guessing that dropouts are introducing randomization, i'll try without that.



Jason Brownlee June 20, 2020 at 6:08 am #

Yes, also there are layers that inject noise directly.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Riccardo June 22, 2020 at 6:09 pm #

i'm sorry if i seem annoying but i do not see layers that can introduce directly noise in the code you provided.

I see convs, batchnorms and concatenates.

Can you please tell me which layer is introducing noise directly?

i think that i have missed something about these layers but reading through the documentation it seems like i know them pretty good.

Start Machine Learning

I really need to understand the position of these noise generator and remove them in order to use a GAN for my application (maybe it could be impossible but i wish to try =))



Jason Brownlee June 23, 2020 at 6:16 am #

Sorry, my mistake, I was thinking of a different GAN.

Sahil Singla June 22, 2020 at 7:17 am #

REPLY ↩

Thanks for the great tutorial.

I have one small doubt:

Do we traverse over the complete dataset?

We passed our entire dataset to the generate_real_samples function and everytime it choose traverse again and again.

So, we might not be traversing over the complete dataset in single epoch?

Please let me know your thoughts.

Thanks.

Start Machine Learning

X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee June 22, 2020 at 1:25 pm #

REPLY ↩

You're welcome.

Correct. On average we cover the whole dataset many times.

Sahil Singla June 22, 2020 at 6:20 pm #

REPLY ↩

So, there is a possibility of missing certain datapoints. This can become a pr

Start Machine Learning

So should I change the code to make sure, it traverse over entire data, or is it still ok, if we don't do that ?



Jason Brownlee June 23, 2020 at 6:17 am #

REPLY ↗

If you prefer. I'm not convinced it makes a difference, but could be a fun experiment.

Riccardo June 23, 2020 at 6:08 pm #

REPLY ↗

oh, ok no problem!

i think that i will investigate stochasticityrought the different convs and batch norm in ord
X input.

best regards

yacine June 30, 2020 at 7:30 pm #

Thank you so much for this super clear explanation and code.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee July 1, 2020 at 5:52 am #

REPLY ↗

You're welcome!

Steve Newbold July 3, 2020 at 11:59 pm #

REPLY ↗

If I wanted to use an input with three colour channels and a target of four colour channels, can this be configured or is it best to just create an additional black 4th channel on the input?

Start Machine Learning

I noticed some greyscale-to-colour models just use the same data in each channel to represent grey images so presumed it must be easier to do this than make the model work with differing numbers of channels.

Also, thanks for the excellent resource!



Jason Brownlee July 4, 2020 at 6:01 am #

REPLY ↗

Off the cuff I recall the images have to have the same number of channels. Perhaps experiment/research and see if you can deviate from this norm.

M July 8, 2020 at 10:42 am #

Thanks for this great post!

For your generator's loss, how can I know if are you minimizing 1: $\log(1 - D(G(x)))$ or maximizing?

How can one change the loss function, any reading suggestions?

Some people say the choice of generator's loss can help the model to not get stuck in early s

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee July 8, 2020 at 1:43 pm #

Good question, this may help you understand the loss function for GANs:

<https://machinelearningmastery.com/how-to-code-the-generative-adversarial-network-training-algorithm-and-loss-functions/>

You can see examples of alternate loss functions in tutorials here:

<https://machinelearningmastery.com/start-here/#gans>

Rao July 21, 2020 at 7:02 am #

REPLY ↗

Hello Jason,

Start Machine Learning

What would be the optimal loss values (Generator and Discriminator loss) of a successful conditional GAN model? Are the values same as an unconditional GAN ? (i.e around 0.7 or 0.6, as mentioned in your unconditional GAN article)

Secondly, I have done the training of pix2pix for a certain image to image translation task in two different ways.

1st method: Trained the discriminator patch outcome against a matrix of real or fake labels (as mentioned in this article)

2nd method: The discriminator still gives a patch, but this time, the patch average was taken and was trained against a single value (i.e avg value of the patch against a real or fake label).

During the training (towards the saving of a good model), the first method, yields a patch avg value of about 0.4 for a real image pair and about 0.3 for a fake image pair.

But the second model, yields a patch avg value of about 0.0004 for both real and fake image pairs.

Both these models yielded a good quality image with its Generator and the Discriminator loss standing around 0.7 and 0.6 respectively. My doubt is why such discrepancy with the avg patch values even though both the models yields a good output. This doesn't make sense even though this model yielded a good translated image.(Because as far as I know, the avg patch value for a real pair should be close to 1 for a real pair and 0 for a fake image pair. This would mean a value of 0.4 for a real pair and 0 for a fake image pair).

What should be the avg patch values for a good model? Any amount of insights into this would be helpful.

Thanks!



Jason Brownlee July 21, 2020 at 1:43 pm #

GANs don't converge, so there is no optimal loss values, you can learn more here:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-gan-not-converging>

Comparing loss between models/runs is not valid.

GANs are hard to evaluate, subjective image quality is about the best we can do, although there are some metrics described here that might help:

<https://machinelearningmastery.com/how-to-evaluate-generative-adversarial-networks/>

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Gruhit Patel July 22, 2020 at 4:56 pm #

Start Machine Learning

Sir, Why exactly are we merging two images in discriminator ?? What effect does it have ?? And why are we not keeping just the colored image in Discriminator ??



Jason Brownlee July 23, 2020 at 6:03 am #

REPLY ↗

The discriminator is given the input image and a target image and comments on whether the target is a real translation or a generated translation.

Jamal July 24, 2020 at 7:54 pm #

Can PIX 2 PIX GAN works for gray-scale images??

what if if we use the above same architecture for gray-scale source and target images



Jason Brownlee July 25, 2020 at 6:17 am #

Modification of the model architecture is required.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

REPLY ↗

Zoya July 24, 2020 at 5:10 pm #

I used different data of source and target image. My source and target images are gray scale.

But when i run the code , the discriminator loss is going to zero with very few iterations but generator loss is very high that is „9782.150 up to so on.

It cannot be decreasingWhat can i do ??



Jason Brownlee July 25, 2020 at 6:13 am #

REPLY ↗

Start Machine Learning

GAN loss does not converge, you can learn more here:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-gan-not-converging>

John July 24, 2020 at 7:44 pm #

REPLY ↗

I have different source and Target images. And My source and target images are in gray scale but my discriminator loss is going to very low reaches to zero but generator loss is very high.

what can I do now ?? Can Pix to Pix GAN work for gray scale images.



Jason Brownlee July 25, 2020 at 6:17 am #

You may need to tune the model – explore – in order to discover how to best model images.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee July 26, 2020 at 1:40 pm #

REPLY ↗

It is hard to know – experimentation is required, perhaps start with tuning the learning rate with a similar network structure adjusted for the changed number of channels.

Jeremy Bolton July 28, 2020 at 7:50 pm #

REPLY ↗

Hi!

Start Machine Learning

Thanks for your great work!

I found your model above and tried with the cityscapes images. I trained ~3000 image pairs from segmentation to photographic pictures. First I convert the images to 256×256 and kept the 100 epochs, then trained with 250 epochs. The results were good, but blurry, so I converted the original 1024×2048 resolution images to 512×512 and trained them till 250 epochs.

The results didn't really improve, but somehow I'd like to get less blurry pictures. I think increasing the number of epochs or the image resolution didn't change a lot, so my question would be: Do I need to change on the architecture of the models? If yes, can you give me a hint what further layers should I use?

Thank you very much and keep up the good work!



Jason Brownlee July 29, 2020 at 5:50 am #

You may need to experiment with the model architecture and learning hyperparameters specific dataset.

Jeremy Bolton July 30, 2020 at 8:40 pm #

Thanks for your reply, Jason.

Can you give me a hint, what architectural changes I should start with if I want to train instead of 256×256? More conv2d layers, dropout layers or multiple discriminators/generators as in pix2pixhd?

Thank you.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee July 31, 2020 at 6:18 am #

REPLY ↗

A good approach is to scale the number of blocks up or down from the current to meet the desired image dimensions.

I would encourage you to experiment and observe the effects on input/output shapes to get a feeling for it.

Start Machine Learning

REPLY ↗

phaneeshwar August 2, 2020 at 4:18 am #

```
dataset = LoadRealData ('C:/Users/Eeshwar/Desktop/deep learning/maps11.npz')
print('Loaded',dataset[0].shape, dataset[1].shape)
Imgshape = dataset[0].shape[1:]
dmodel = DiscriminatorModel(Imgshape)
gmodel = GeneratorModel(Imgshape)
ganmodel = GANModel(dmodel,gmodel,Imgshape)

TrainModel(dmodel,gmodel, ganmodel,dataset)
```

Loaded (1096, 256, 256, 3) (1096, 256, 256, 3)

WARNING:tensorflow:Discrepancy between trainable weights and collected trainable weights

model.compile after ?

InvalidArgumentException: data[0].shape = [4] does not start with indices[0].shape = [2]

[[{{node training/Adam/gradients/gradients/loss_3/dense_2_loss/Mean_grad/DynamicStitch}}]]

Sir could you please help me to resolve this issue. I Thank You in advance

**Jason Brownlee** August 2, 2020 at 5:47 am #

Sorry to hear that, this will help:

<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me>

did you set model._trainable without calling

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Dhruv Agarwal August 2, 2020 at 4:41 pm #

REPLY ↗

Hello sir, the tutorial was great, but i have 2 questions.

1) In the define_discriminator() function, you have set the loss_weights parameter to 0.5, to slow down the training of discriminator. Can't we reduce the learning rate of the discriminator model to slow the training, instead of specifying the loss_weights parameter?

2) In the define_gan() function, why was there even a need to specify loss_weights parameter over there?

Start Machine Learning



Jason Brownlee August 3, 2020 at 5:45 am #

REPLY ↗

Thanks.

Perhaps try it and see.

We do see a loss_weights for the gan.

Dhruv Agarwal August 4, 2020 at 1:14 am #

REPLY ↗

Ok , i will try reducing the learning rate instead of specifying the loss_weights parameter in the define_discriminator() . But i am sorry, but i still do not get the answer of the second question, i.e, why do we need to specify a loss function.



Jason Brownlee August 4, 2020 at 6:42 am #

To match the implementation described in the paper.

It has the effect of giving most attention to L1 and a tiny bit of attention to cross entropy.

This is explained in that section of the tutorial, perhaps re-read?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Amine Zera August 4, 2020 at 9:35 pm #

REPLY ↗

Hello sir, thank you for the great tutorial !

I am new to Machine Learning,

I want to change the clothings of people in images or videos. So i should train pix2pix on a clothes dataset ?

The second question is that i dont want to change anything else in the image except the clothes, so if i apply pix2pix on the image it will change everything, how can i target only clothing in a image ?

Thank's again for your great work !

Start Machine Learning



Jason Brownlee August 5, 2020 at 6:11 am #

REPLY ↗

Perhaps try it and see how well it can do.

Alex Westcott August 19, 2020 at 11:04 pm #

REPLY ↗

Hi,

I have trained the exact model outlined in the tutorial with the same data-set quite a few times and the losses of the discriminator are always consistently 0.000 after around 5000 steps. Looking at the loss to more significant figures, you state that, if the discriminator loss stays at zero for a long time then there is training failure.

The generator still improves after the discriminator loss states 0.000, however I presume that impact on the training of the generator.

Thank you for the great tutorial, it helped a lot!



Jason Brownlee August 20, 2020 at 6:43 am #

X

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Zero loss indicates a failure mode:

<https://machinelearningmastery.com/practical-guide-to-gan-failure-modes/>

Recall that GANs do not converge:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-gan-not-converging>

Are you saving models along the way during training?

Are you able to inspect the progress of training, does it get good then go bad or is it bad the entire time?

Alex Westcott August 20, 2020 at 6:02 pm #

REPLY ↗

Start Machine Learning

I am saving the model every 5 epochs, and the predicted images do improve slightly during training, and by the end look reasonably good, (I presume that the discriminator hasn't had an impact on the quality and it is just the generator improving by itself). The losses of both the discriminator and generator decrease to start with, but the discriminator slowly decreases to 0 and the generator stays pretty low (between 1 and 5).

I have assumed that the discriminator is too good at determining the real and fake images, as I have removed a few layers from it and its loss doesn't decay to 0 during training.



Jason Brownlee August 21, 2020 at 6:26 am #

REPLY ↗

Interesting. Thanks for sharing Alex.

Reshma Jindal August 20, 2020 at 1:51 am #

I have around 3700 images to train on.

Can you roughly guide for the hyperparameters(like n_epochs,n_batch to be set as I'm encountering error while training).
Please help in resolving it.

/home/reshmajindal/.local/lib/python3.6/site-packages/keras/engine/training.py:490: UserWarning: Discrepancy between trainable weights and collected trainable weights, did you set `model.trainable` without calling `model.compile` after changing the model's structure?
'Discrepancy between trainable weights and collected trainable'

Killed

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Reshma Jindal August 20, 2020 at 1:53 am #

REPLY ↗

<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me/>
I even tried all of this, but didn't work.

Start Machine Learning



Jason Brownlee August 20, 2020 at 6:49 am #

REPLY ↗

We cannot know the best way to configure the model, instead we must use experiments to tune and discover what configuration works best for a given dataset.

These tutorials will teach you:

<https://machinelearningmastery.com/start-here/#gans>

Awadelrahman M. A. Ahmed August 31, 2020 at 2:33 am #

REPLY ↗

Thanks for this GREAT detailed tutorial. One question I have in mind is how to adapt training/validation images have different height and width values?



Jason Brownlee August 31, 2020 at 6:17 am #

You're welcome.

Typically images are all resized to the same width and height expected by the model.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

REPLY ↗

Awadelrahman M. A. Ahmed September 2, 2020 at 6:17 am #

resizing is a bit flexible term 😊

cropping big images leads to loosing some information. enlarging small images might lead to blurry images. Super-resolution is computationally expensive and needs auxiliary models. What do you think the good way to "resize" images to work properly with this model ?



Jason Brownlee September 2, 2020 at 6:35 am #

REPLY ↗

Start Machine Learning

I recommend exploring many different approaches and discover what works best for your specific project.

Awadelrahman M. A. Ahmed September 2, 2020 at 9:02 am #

YES!! the best way to find out is by doing it! this why I feel addicted to this machinelearningmastery :p



Jason Brownlee September 2, 2020 at 1:29 pm #

Thanks!

Nalin Nagar September 9, 2020 at 11:36 am #

Is there a way to input your own image? I haven't seen any demonstrations that are myself but to no avail.



Jason Brownlee September 9, 2020 at 1:34 pm #

Yes, the last part of the tutorial shows this.

Harry September 9, 2020 at 4:10 pm #

Hi everyone. Thank you for super guideline for implementation. I have one question. Can i generate 1024x1024 px image by using pix2pix-GAN?

REPLY ↗

Start Machine Learning



Jason Brownlee September 10, 2020 at 6:22 am #

REPLY ↗

Perhaps try scaling up the model for large images and see what kind of results you get.

I would expect quality to fall off. It might be easier with a model based on the progressive-growing architecture.

Harry September 9, 2020 at 4:23 pm #

REPLY ↗

By the way, my dataset image size is smaller than 1024px

Bidesh Sengupta September 14, 2020 at 2:52 pm #

Hi!

It is a really good tutorial. I wish to apply this concept to my work. But I want to give some number of images as input and wish to get the image as output.

Can you guide me on how to change the code to implement this? Is it at all possible?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee September 15, 2020 at 5:17 am #

REPLY ↗

Thanks.

Perhaps take a look at some alternate GANs like conditional GAN or InfoGAN:

<https://machinelearningmastery.com/how-to-develop-an-information-maximizing-generative-adversarial-network-infogan-in-keras/>

Manohar Sai October 5, 2020 at 10:56 pm #

REPLY ↗

Start Machine Learning

Thanks for this great tutorial.

Both losses for the discriminator has gone to zero in the first 100 epochs.

Can you help me?



Jason Brownlee October 6, 2020 at 6:51 am #

REPLY ↗

Perhaps restart training and stop once the generated images are good enough.

Manohar October 5, 2020 at 11:18 pm #

REPLY ↗

Great tutorial sir.

I have my both discriminator loss heading to zero, in the first 200 steps. I cannot solve my issue with the version?



Jason Brownlee October 6, 2020 at 6:52 am #

X

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Perhaps try changing/tuning the learning hyperparameters of the model.

Perhaps try some of the suggestions here:

<https://machinelearningmastery.com/how-to-code-generative-adversarial-network-hacks/>

Manohar October 6, 2020 at 7:17 pm #

REPLY ↗

Thanks for the reply sir. I will try that.

But why shouldn't it work if the code is same as above?

Start Machine Learning



Jason Brownlee October 7, 2020 at 6:43 am #

REPLY ↗

The model uses a stochastic learning algorithm, you may need to run the example a few times to get a good result.

you can learn more about this here:

<https://machinelearningmastery.com/faq/single-faq/why-do-i-get-different-results-each-time-i-run-the-code>

Mick October 28, 2020 at 8:06 pm #

REPLY ↗

Great tutorial!

I am trying to apply this architecture to a MRI image-to-image translation task. I have two questions:

1) After slicing the MRI data to 2D slices. Do I need to convert the NIFTI-files to JPEG or can I work with numpy arrays?

2) MRI images are grayscale whereas the example code in this tutorial uses RGB images. Will I need to change anything to deal with grayscale images?

Thanks Jason.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee October 29, 2020 at 8:00 am #

REPLY ↗

The model takes image data as numpy arrays. I don't know if converting data to jpeg first is required for your data.

Yes, you can adapt the model for grayscale images, e.g. change the number of channels for input images to D() and output from G().

Adrien November 2, 2020 at 7:08 am #

REPLY ↗

Jason,

What an incredible article. I reproduced your methodology on a research project on mechanical networks, where the model learns to draw mechanical linkages between parts of the system. It works perfectly, despite a small sample size.

Start Machine Learning

I re-used one of your images (the Unet architecture of the Generator) on a blog post I made on Medium, carefully citing your article as source and your work as reference. You can check it out here:

<https://adriensaremi.medium.com/develop-a-image-to-image-translation-model-to-capture-local-interactions-in-mechanical-networks-9c2f45230849>

I wanted to make sure you approved the re-use of the image in question. Thanks again for your work here and more broadly on Machine Learning Mastery.



Jason Brownlee November 2, 2020 at 7:52 am #

REPLY ↗

Thanks.

That's fine, well done on your post!

ZiZi November 11, 2020 at 2:13 am #

Thank you for your great tutorial

I read a few posts about GANs and I realized GANs applyed in square images. is it right? can

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee November 11, 2020 at 6:50 am #

Yes, they can, bit square is much simpler.

Rekka Mastouri December 1, 2020 at 11:04 pm #

REPLY ↗

Thank you for your great tutorial.

please how can I use GAN for deformable image registration?

Thanks Jason.

Start Machine Learning



Jason Brownlee December 2, 2020 at 7:44 am #

REPLY ↗

You're welcome.

Perhaps start by checking the literature for existing approaches and try them first.

Zaineb December 2, 2020 at 11:15 pm #

REPLY ↗

Hi,

Hope you are doing good.

I have tried your code and it works perfectly well.

I need to know, how about testing this module on a separate dataset,because i have found out that i need to include testing dataset also.

If i use a part of validation dataset (and call it my test dataset) on saved model (e.g model_1) to test this test dataset, the segmentation results are not desirable.

I would be glad if you can shed some light on this. Also please tell me, is there any way that the pix2pix model there any reference that signifies that testing pix2pix for image to image translation is not a good idea?

Thanks

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee December 3, 2020 at 8:18 am #

REPLY ↗

Thanks.

Sorry, I don't have an example of combining GAN output with a predictive model – I don't think I can give you good off the cuff advice on the topic. Perhaps check the literature.

Michal December 16, 2020 at 12:15 pm #

REPLY ↗

Start Machine Learning

Hey, really well explained, good job!

I have implemented similar cGAN for b&w image colorization. It is very hard to train, and somehow after many, many epochs on big datasets I got some 'good enough' results, but I wonder how can I measure accuracy for translated images?

Also during training and after finishing it my cGAN is resulting in very big Losses of gen like 10.0 and 2.0 at the end of training. Disctiminator's loss is near 0 and peaking sometimes to even 3 or 5. How can I measure accuracy of trained model or during training?

Thanks



Jason Brownlee December 16, 2020 at 1:41 pm #

REPLY ↗

Thanks.

Good question, this may give you some ideas of how to evaluate a GAN model:

<https://machinelearningmastery.com/how-to-evaluate-generative-adversarial-networks/>

Muhammad Gulfam December 27, 2020 at 7:34 am #

Hi Jason,

Thank you very much for detailed explanation with examples. It is very helpful.

I am trying to edit the code through notepad++ but it is giving me indentation error. Seems like

Can you please tell me what IDE or editor you used?

Apologies for a silly question.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee December 27, 2020 at 9:26 am #

REPLY ↗

You're welcome.

This will help you copy the code correctly:

<https://machinelearningmastery.com/faq/single-faq/how-do-i-copy-code-from-a-tutorial>

Start Machine Learning

Gavin December 28, 2020 at 6:01 am #

REPLY ↗

Amazing tutorial, even more impressive that you've responded to every comment over year later! Quick question: you said that if either discriminant loss plateaus at 0 for an extended period of time that it has most likely failed and should be restarted. I am running it for the third time and both have landed on zero again, am I doing something wrong? Anything I can do to improve chances of it succeeding or just keep trying? (P.S. I am using different images and am using 2000 images as opposed to your ~1100 (still 100 epochs) but I assume that this does not affect the base of the model). Thanks in advance.



Jason Brownlee December 28, 2020 at 6:06 am #

REPLY ↗

Thanks!

Sorry to hear that.

Perhaps try fewer epochs?

Perhaps try changing other learning hyperparameters?

Perhaps try adjusting the architecture?

Perhaps try some of the ideas here:

<https://machinelearningmastery.com/how-to-code-generative-adversarial-network-hacks/>

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Gavin December 30, 2020 at 10:34 am #

Just an update, tried some things you suggested, in the article as well as in the well appreciated comment, not much changed. Let it run just to see what would happen and even though the model read 0 for both discriminators for more than 6 epochs, it still gave me decent results, so I'm happy. Thanks for the amazing article and the helpful advice, will definitely be reading up on some of your other articles.



Jason Brownlee December 30, 2020 at 10:44 am #

REPLY ↗

Thanks for the update and well done!

Start Machine Learning

Eric January 2, 2021 at 4:51 am #

REPLY ↗

Hi, I love your whole blog and tutorials!

Just a question, Is it possible to train a model that uses 2 source images for one target?

For example from a traditional satellite image + an Infra Red (IR) image recreate the corresponding map?

Thanks a lot



Jason Brownlee January 2, 2021 at 6:27 am #

REPLY ↗

Thanks!

I don't see why not. I expect there are papers on exactly this – I recommend seeking them

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee January 3, 2021 at 1:29 pm #

REPLY ↗

You're welcome.

Yes, you can load the saved model and continue training. You can use the same code as the first round of training as a starting point.

D.GHOSH January 11, 2021 at 6:02 am #

REPLY ↗

Start Machine Learning

Is this model applicable to generate super resolution data?



Jason Brownlee January 11, 2021 at 6:23 am #

REPLY ↗

No, I believe there are more specialized models for that problem described in the literature.

Muhammad Gulfam January 12, 2021 at 8:38 am #

REPLY ↗

Can you please share the link of some articles for those specialized models? for generating super resolution data



Jason Brownlee January 12, 2021 at 10:33 am #

You can search for papers on the topic here:

<https://scholar.google.com/>

Muhammad Gulfam January 12, 2021 at 8:06 am #

What is the significance of converting the pixel values from [0, 255] to [-1, 1]?

Is it because of the tanh activation function being used in the generator model for the last layer?

This architecture can be used to matrix to matrix mapping as well. but a matrix might have pixel (arr[row, col]) values as real values (from [0, inf] instead of [0, 255]). In that case, what would you suggest for transformation (to [-1, 1])? Should that still be done?

Thanks,

Apologies for multiple questions.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee January 12, 2021 at 10:31 am #

Start Machine Learning

Yes, exactly.

Yes, it is standard practice to use tanh for the output layer of gan generator models and to scale data to match the distribution of the activation function.

Muhammad Gulfam January 12, 2021 at 5:05 pm #

REPLY ↗

Thank you very much. I appreciate the responses.



Jason Brownlee January 13, 2021 at 6:10 am #

You're welcome.

Muhammad Gulfam January 15, 2021 at 5:11 am #

I have noticed that in the code that the discriminator model is being compiled and the model is not being compiled. generator is being saved. Whenever I load the generator model "No training configuration found in save file: the model was *not* compiled. Compile it manual. Can you please guide me if it can affect model's performance? seems like my models are not After googling it I got a perception that it is just a warning but still wanted to check with you.

Thanks

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee January 15, 2021 at 5:59 am #

REPLY ↗

No need as we are not training it directly. You can ignore the warning.

Start Machine Learning

Muhammad Gulfam January 15, 2021 at 1:25 pm #

REPLY ↗

Thank you for the response.

WinK January 21, 2021 at 2:43 pm #

REPLY ↗

I am fan of your site. Always thanks for your great article.

I would like to ask loss function that you utilize in the logical gan model. In your code block, 2 loss function was used in define_gan function.

```
model.compile(loss=['binary_crossentropy', 'mae'], optimizer=opt, loss_weights=[1,100])
```

If I understand correctly, 'mae' takes labels (true and predicted labels) instead of images. But

$$L1(G) = \mathbb{E}_{x,y,z} [\|y - G(x,z)\|]$$

The output of G model is image and their loss function is defined based on differences between labels.

Is it the same effect with labels instead of using images?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



Jason Brownlee January 22, 2021 at 7:15 am #

Yes, MAE is the L1 norm between image pixels.

WinK January 25, 2021 at 12:54 pm #

REPLY ↗

Thank you for your answer.

Muhammad Gulfam February 1, 2021 at 12:02 pm #

Start Machine Learning

Hi Dr. Brownlee,

In your last version there was a line in the define_gan method:

```
# make weights in the discriminator not trainable
```

```
d_model.trainable = False
```

my question is that if the discriminator is not trainable then how will it improve?

In current version of your code you have replaced it by following lines:

```
# make weights in the discriminator not trainable
```

```
for layer in d_model.layers:
```

```
if not isinstance(layer, BatchNormalization):
```

```
layer.trainable = False
```

if the weights are not trainable then how will discriminator learn and get better, and contribute to make the generator better?

my understanding was that weights are the ones that are supposed to be trained in the training process. Please correct me if I am wrong.

Apologies as I am not an expert. I am learning through your articles and other stuff.

Thanks in advance.



Jason Brownlee February 1, 2021 at 1:48 pm #

The D is only not trainable when part of the composite model. This is called layer model.

Start Machine Learning X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

REPLY ↩

Muhammad Gulfam February 2, 2021 at 2:30 am #

Thank you so much Dr. Brownlee.

REPLY ↩

Glenn Q February 16, 2021 at 10:23 am #

Hi Dr. Brownlee, if I want to have a higher learning rate for the discriminator and a lower one for the generator, say $2e^{-4}$ for discriminator and $1e^{-4}$ for the generator, should I just change the learning rate setting of the composite model?

Start Machine Learning



Jason Brownlee February 16, 2021 at 1:38 pm #

REPLY ↗

Yes, the composite model is used to update the generator.

Let me know how you go with your approach.

Alice February 17, 2021 at 8:19 pm #

REPLY ↗

Hi Jason,

Thank you for your great tutorial.

I just want to ask you one question: why during the inference we have to keep the batch norm?

I understand that the dropout is performed to add some noise, but I thought it was necessary

Moreover, I have performed the training with a batch size = 1 and in the prediction phase I have images of dimension [N, 256, 256, 3] and the results were very different. Using a batch size = 1 think that this is correlated to the adoption of BN in training modality.

Thank you for your time

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE



Jason Brownlee February 18, 2021 at 5:14 am #

No. Batchnorm and dropout can be flipped to inference mode. Batchnorm will use learned mu and sigma and dropout will stop dropping out.

Alice February 23, 2021 at 10:16 pm #

REPLY ↗

but the flag training is set to True for both BN and dropout, I think that this flag makes them work as during the training phase

Start Machine Learning

[REPLY ↗](#)**anarchitect** February 28, 2021 at 5:21 am #

Hi,

Is it possible to plot losses in realtime? I couldn't manage to do it. Could you please help me?

[REPLY ↗](#)**Jason Brownlee** February 28, 2021 at 5:42 am #

Perhaps via tensorboard?

Lin April 10, 2021 at 4:55 pm #

Hello, thank you for the sharing.

I'd like to know what is d1[0.362] d2[0.405] g[78.143] each loss value's meaning?

Does it mean that is fake when discriminator's loss value close to zero?

And what is the composite's loss value mean?

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)[X](#)**Jason Brownlee** April 11, 2021 at 4:49 am #

The values are hard to interpret.

Nevertheless d1 and d2 are the discriminators loss on real samples and fake samples (of "B") respectively, and g is loss of the composite model on real samples.

This may help you interpret cross entropy more generally:

<https://machinelearningmastery.com/cross-entropy-for-machine-learning/>

Nada April 27, 2021 at 1:39 pm #[Start Machine Learning](#)

Hi Jason,

Thank you very much for this informative article.

I have a question about a good GAN model to create more synthesis images from a small set of medical images? Is styleGaN good for this problem?



Jason Brownlee April 28, 2021 at 6:00 am #

REPLY ↗

Perhaps trial a few methods and discover what works well or best for your dataset.

Rishabh singh April 27, 2021 at 11:53 pm #

Hey,

If possible , can you please share the .h5 model after complete training. As I am trying but not getting power.

I have tried on colab too, but gets stopped after some time.



Jason Brownlee April 28, 2021 at 6:02 am #

Sorry, I cannot share saved models.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

START MY EMAIL COURSE

Leave a Reply

Start Machine Learning

Name (required) Email (will not be published) (required) Website

SUBMIT COMMENT

**Welcome!**

I'm Jason Brownlee PhD

and I **help developers** get results with **machine learning**.

[Read more](#)

Never miss a tutorial:**Picked for you:**

[How to Develop a Pix2Pix GAN for Image-to-Image Translation](#)

<https://machinelearningmastery.com/how-to-develop-a-pix2pix-gan-for-image-to-image-translation/>

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

 Email Address[START MY EMAIL COURSE](#)

X

[Start Machine Learning](#)



How to Develop a 1D Generative Adversarial Network From Scratch in Keras

How to Develop a CycleGAN for Image-to-Image Translation with Keras

How to Develop a Conditional GAN (cGAN) From Scratch

How to Train a Progressive Growing GAN in Keras for Synthesizing Faces

Start Machine Learning X

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

[START MY EMAIL COURSE](#)

Loving the Tutorials?

The **GANs with Python** EBook is
where you'll find the **Really Good** stuff.

[>> SEE WHAT'S INSIDE](#)

Start Machine Learning ×

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning